
Network Applications: High-performance Server Design

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

10/21/2021

Outline

- ❑ Admin and recap
- ❑ High-performance network server design

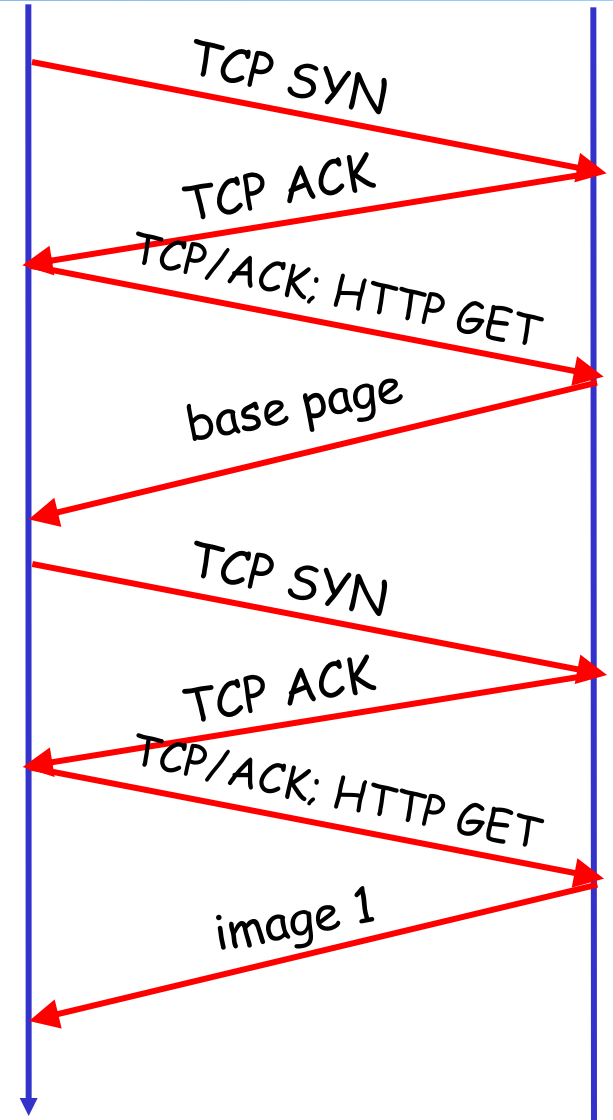
Admin

- ❑ Lab Assignment Three
 - Part 1: Due Nov. 11
 - Part 2: To be posted

- ❑ Exam 1 date?

Recap: Latency of Basic HTTP/1.0

- ≥ 2 RTTs per object:
 - TCP handshake --- 1 RTT
 - client request and server responds --- at least 1 RTT (if object can be contained in one packet)

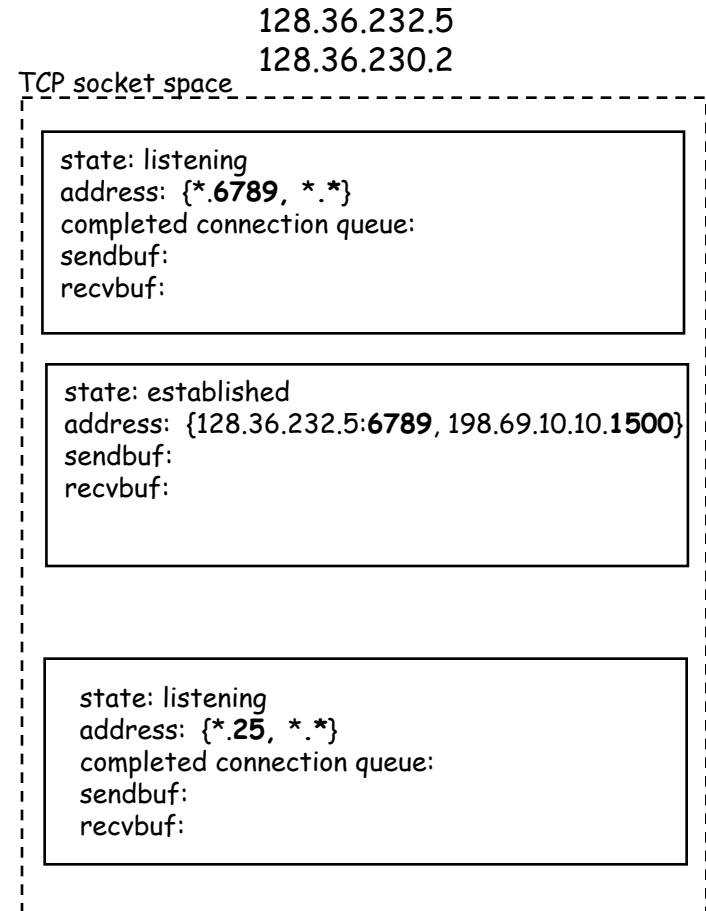
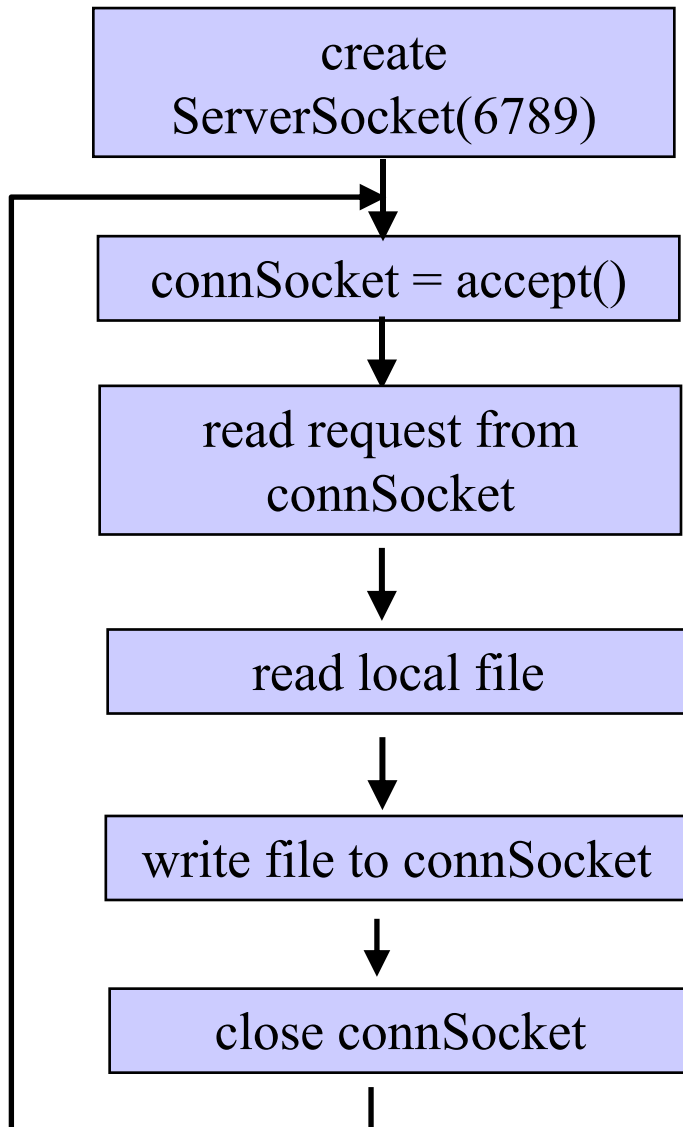


Recap: Substantial Efforts to Speedup HTTP/1.0

- ❑ Reduce the number of objects fetched [Browser cache]
- ❑ Reduce data volume [Compression of data]
- ❑ Header compression [HTTP/2]
- ❑ Reduce the latency to the server to fetch the content [Proxy cache]
- ❑ Remove the extra RTTs to fetch an object [Persistent HTTP, aka HTTP/1.1]
- ❑ Increase concurrency [Multiple TCP connections]
- ❑ Asynchronous fetch (multiple streams) using a single TCP [HTTP/2]
- ❑ Server push [HTTP/2]



WebServer Implementation

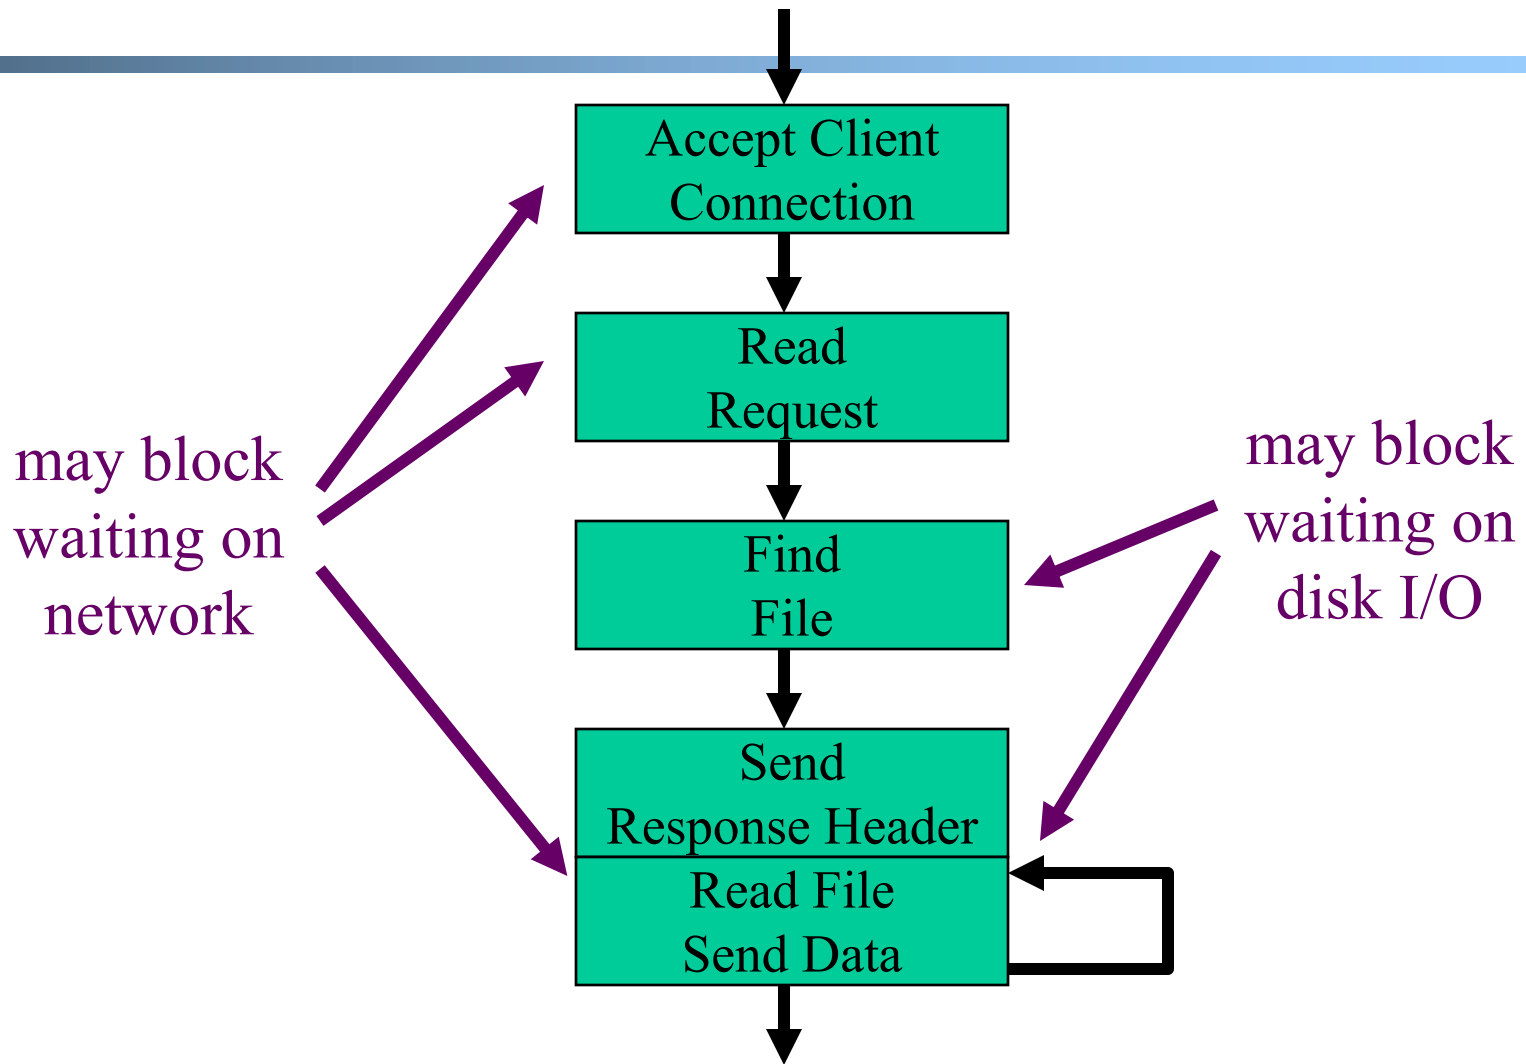


Discussion: what does each step do and how long does it take?

Demo

- ❑ Try TCPServer
- ❑ Start two TCPClient
 - Client 1 starts early but stops
 - Client 2 starts later but inputs first



Server Processing Steps



Writing High Performance Servers: Major Issues

- ❑ Many socket and IO operations can cause a process to block, e.g.,
 - `accept`: waiting for new connection;
 - `read` a socket waiting for data or close;
 - `write` a socket waiting for buffer space;
 - `I/O read/write` for disk to finish

Goal: Limited Only by Resource Bottleneck



Outline

- ❑ Admin and recap
- ❑ Network server design
 - Overview
 - *Multi-thread network servers*


Multi-Threaded Servers

❑ Motivation:

- Avoid blocking the whole program (so that we can reach bottleneck throughput)

❑ Idea: introduce threads

- A thread is a sequence of instructions which may execute in parallel with other threads
- When a blocking operation happens, only the flow (thread) performing the operation is blocked




Background: Java Thread Model


- ❑ Every Java application has at least one thread
 - The “main” thread, started by the JVM to run the application's main() method
 - Most JVMs use POSIX threads to implement Java threads

- ❑ main() can create other threads
 - Explicitly, using the Thread class
 - Implicitly, by calling libraries that create threads as a consequence (RMI, AWT/Swing, Applets, etc.)

Thread vs Process



A computer



*Processes do not
share data*

*Threads share data
within a process*

Creating Java Thread

- ❑ Two ways to implement Java thread
 1. Extend the `Thread` class
 - Overwrite the `run()` method of the `Thread` class
 2. Create a class `C` implementing the `Runnable` interface, and create an object of type `C`, then use a `Thread` object to wrap up `C`
- ❑ A thread starts execution after its `start()` method is called, which will start executing the thread's (or the `Runnable` object's) `run()` method
- ❑ A thread terminates when the `run()` method returns

Option 1: Extending Java Thread

```
class PrimeThread extends Thread {  
    long minPrime;  
  
    PrimeThread(long minPrime) {  
        this.minPrime = minPrime;  
    }  
  
    public void run() {  
        // compute primes larger than minPrime ...  
    }  
}  
  
PrimeThread p = new PrimeThread(143);  
p.start();
```


Option 1: Extending Java Thread

```
class RequestHandler extends Thread {  
    RequestHandler(Socket connSocket) {  
        // ...  
    }  
    public void run() {  
        // process request  
    }  
    ...  
}  
  
Thread t = new RequestHandler(connSocket);  
t.start();
```

Option 2: Implement the Runnable Interface

```
class PrimeRun implements Runnable {  
    long minPrime;  
    PrimeRun(long minPrime) {  
        this.minPrime = minPrime;  
    }  
  
    public void run() {  
        // compute primes larger than minPrime . . .  
    }  
}  
  
PrimeRun p = new PrimeRun(143);  
  
new Thread(p).start();
```


Example: a Multi-threaded TCPServer

- ❑ Turn TCPServer into a multithreaded server by creating a thread for each accepted request

Per-Request Thread Server

```
main() {  
    ServerSocket s = new ServerSocket(port);  
    while (true) {  
        Socket conSocket = s.accept();  
        RequestHandler rh  
            = new RequestHandler(conSocket);  
        Thread t = new Thread (rh);  
        t.start();  
    }  
}
```

```
class RequestHandler implements Runnable {  
    RequestHandler(Socket connSocket) { ... }  
    public void run() {  
        //  
    } }  
}
```



Try the per-request-thread TCP server: `TCPServerMT.java`

Summary: Implementing Threads


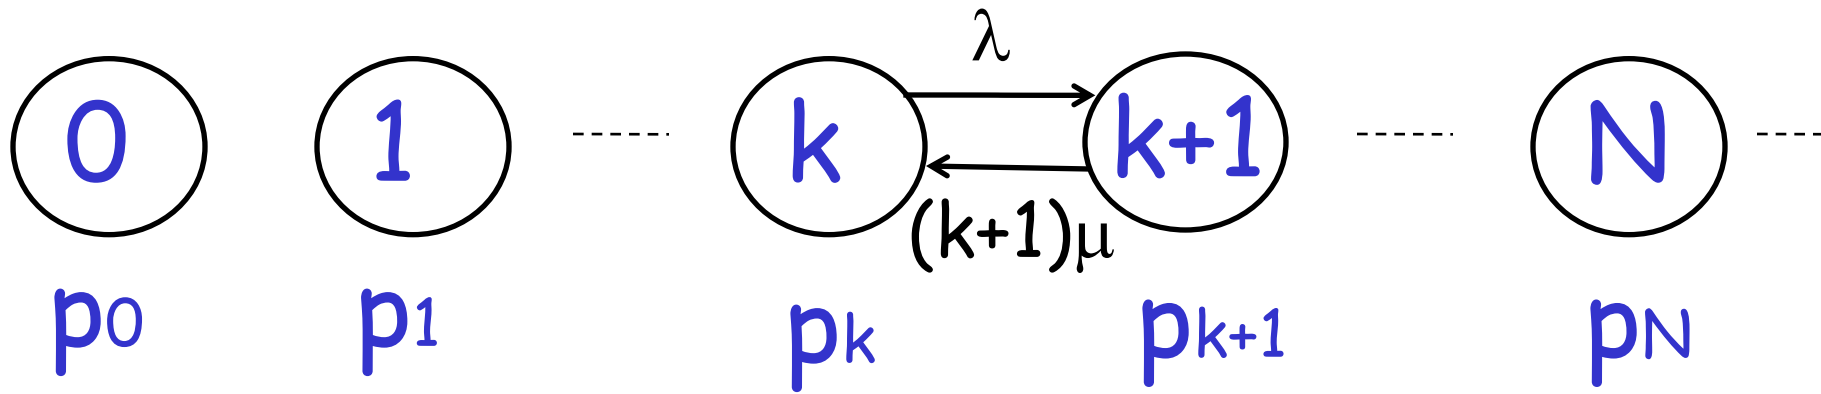
```
class RequestHandler
    extends Thread {
    RequestHandler(Socket connSocket)
    {
        ...
    }
    public void run() {
        // process request
    }
    ...
}

Thread t = new RequestHandler(connSocket);
t.start();
```

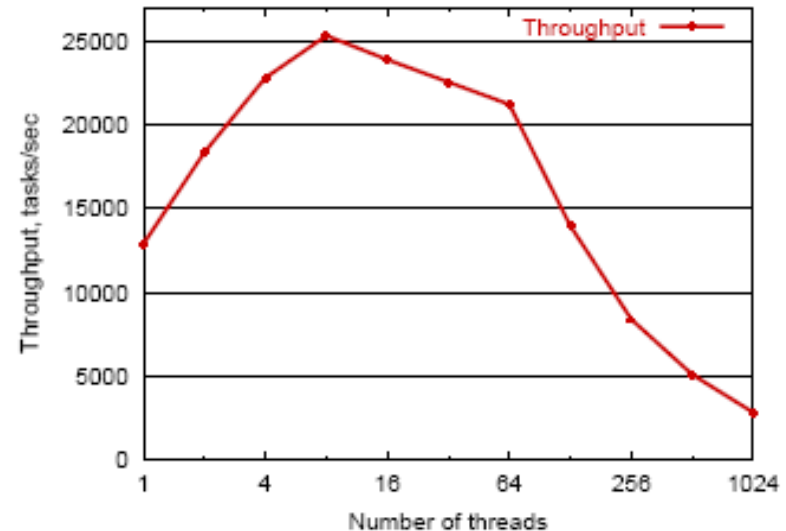
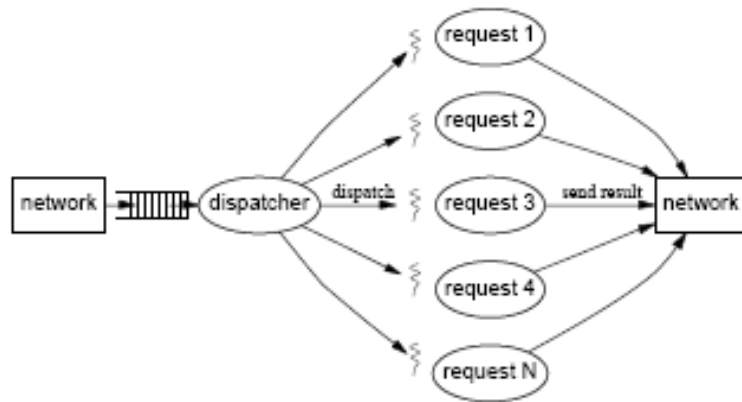
```
class RequestHandler
    implements Runnable {
    RequestHandler(Socket connSocket)
    {
        ...
    }
    public void run() {
        // process request
    }
    ...
}

RequestHandler rh = new
    RequestHandler(connSocket);
Thread t = new Thread(rh);
t.start();
```

Modeling Per-Request Thread Server: Theory



Problem of Per-Request Thread: Reality



(937 MHz x86, Linux 2.2.14, each thread reading 8KB file)

- ❑ High thread creation/deletion overhead
- ❑ Too many threads → resource overuse → throughput meltdown → response time explosion
 - Q: given avg response time and connection arrival rate, how many threads active on avg?


Background: Little's Law (1961)



□ For any system with no or (low) loss.

□ Assume

- mean arrival rate λ , mean time R at system, and mean number Q of requests at system




□ Then relationship between Q , λ , and R :

$$Q = \lambda R$$

Example: XMU admits 3000 students each year, and mean time a student stays is 4 years, how many students are enrolled?



Little's Law: Proof

$$Q = \lambda R$$



$$\lambda = \frac{A}{t} \quad R = \frac{Area}{A} \quad Q = \frac{Area}{t}$$

Discussion: How to Address the Issue




(937 MHz x86, Linux 2.2.14, each thread reading 8KB file)

Outline

- ❑ Admin and recap
 - ❑ High-performance network server design
 - Overview
 - Threaded servers
 - Per-request thread
 - problem: large # of threads and their creations/deletions may let overhead grow out of control
- *Thread pool*

Using a Fixed Set of Threads (Thread Pool)


- ❑ Design issue: how to distribute the requests from the welcome socket to the thread workers



Design 1: Threads Share Access to the welcomeSocket

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = welcomeSocket.accept();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```

sketch; not
working code

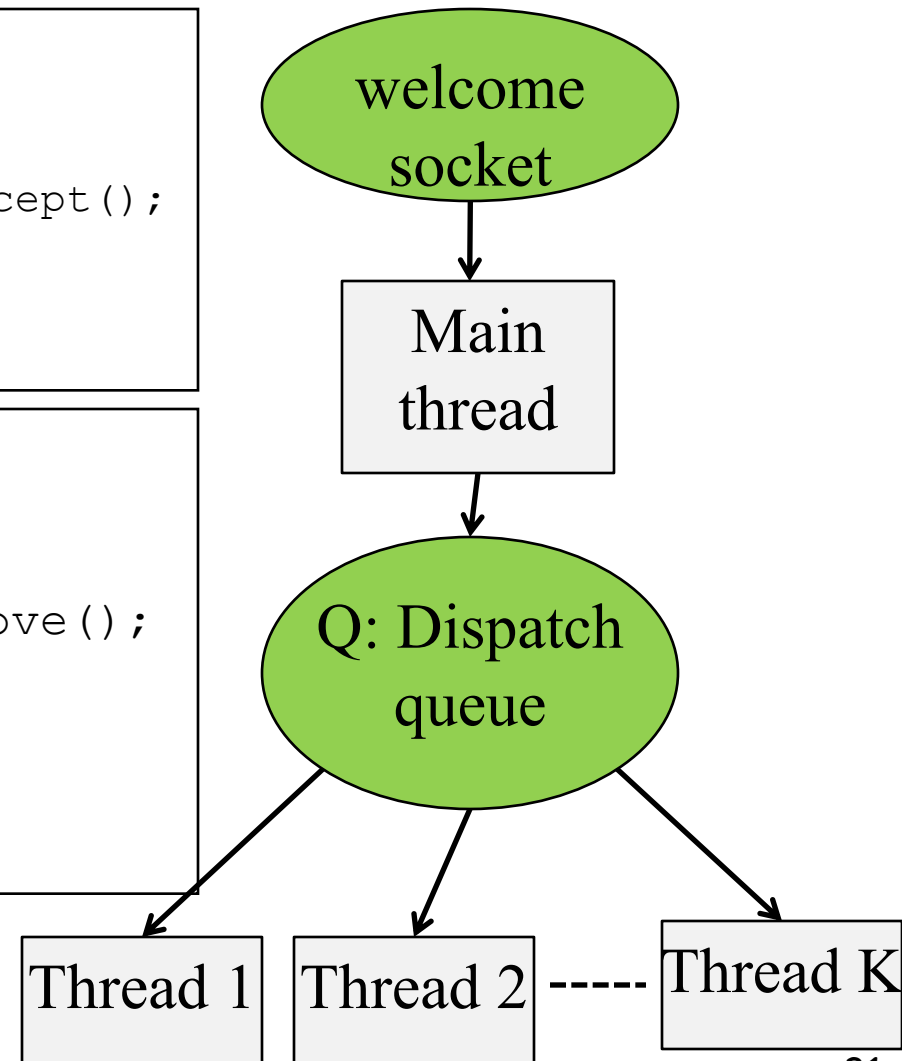


Design 2: Producer/Consumer

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```

sketch; not
working code



Common Issues Facing Designs 1 and 2

- ❑ Both designs involve multiple threads modifying the same data concurrently
 - Design 1: `welcomeSocket`
 - Design 2: `Q`

- ❑ In our original `TCPServerMT`, do we have multiple threads modifying the same data concurrently?

Concurrency and Shared Data

- ❑ Concurrency is easy if threads don't interact
 - Each thread does its own thing, ignoring other threads
 - Typically, however, threads need to communicate/coordinate with each other
 - Communication/coordination among threads is often done by *shared data*

Simple Example

```
public class ShareExample extends Thread {  
    private static int cnt = 0; // shared state, count  
                                // total increases  
  
    public void run() {  
        int y = cnt;  
        cnt = y + 1;  
    }  
  
    public static void main(String args[]) {  
        Thread t1 = new ShareExample();  
        Thread t2 = new ShareExample();  
        t1.start();  
        t2.start();  
        Thread.sleep(1000);  
        System.out.println("cnt = " + cnt);  
    }  
}
```

Q: What is the result of the program?

Simple Example

What if we add a println:

```
int y = cnt;
```

```
System.out.println("Calculating...");
```

```
cnt = y + 1;
```

What Happened?

- ❑ A thread was preempted in the middle of an operation
- ❑ The operations from reading to writing `cnt` should be *atomic* with no interference access to `cnt` from other threads
- ❑ But the scheduler interleaves threads and caused a *race condition*
- ❑ Such bugs can be extremely hard to reproduce, and also hard to debug

Synchronization

- ❑ Refers to mechanisms allowing a programmer to control the execution order of some operations across different threads in a concurrent program.
- ❑ We use Java as an example to see synchronization mechanisms
- ❑ We'll look at locks first.

Java Lock (1.5)

```
interface Lock {  
    void lock();  
    void unlock();  
    ... /* Some more stuff, also */  
}  
class ReentrantLock implements Lock { ... }
```

- ❑ Only one thread can hold a lock at once
- ❑ Other threads that try to acquire it *block (or become suspended)* until the lock becomes available
- ❑ *Reentrant lock can be reacquired by same thread*
 - As many times as desired
 - No other thread may acquire a lock until it has been released the same number of times that it has been acquired
 - Do not worry about the reentrant perspective, consider it a lock

Java Lock

❑ Fixing the ShareExample.java problem

```
import java.util.concurrent.locks.*;
public class ShareExample extends Thread {
    private static int cnt = 0;
    static Lock lock = new ReentrantLock();

    public void run() {
        lock.lock();
        int y = cnt;
        cnt = y + 1;
        lock.unlock();
    }
    ...
}
```

Java Lock

- ❑ It is recommended to use the following pattern



```
...  
lock.lock();  
try {  
    // processing body  
} finally {  
    lock.unlock();  
}  

```

Java synchronized

- ❑ This pattern is really common
 - Acquire lock, do something, release lock after we are done, **under any circumstances, even if exception was raised, the method returned in the middle, etc.**
- ❑ Java has a language construct for this
 - `synchronized (obj) { body }`
 - ❑ Utilize the design that every Java object has its own **implicitly lock** object, also called the **intrinsic lock**, **monitor lock** or simply **monitor**
 - Obtains the lock associated with **obj**
 - Executes **body**
 - Release lock when scope is exited
 - Even in cases of exception or method return

Discussion



- ❑ An object and its associated lock are different !
- ❑ Holding the lock on an object does not affect what you can do with that object in any way
- ❑ Examples:
 - `synchronized(o) { ... } // acquires lock named o`
 - `o.f (); // someone else can call o's methods`
 - `o.x = 3; // someone else can read and write o's fields`

Synchronization on this

```
class C {  
    int cnt;  
    void inc() {  
        synchronized (this) {  
            cnt++;  
        } // end of sync  
    } // end of inc  
}
```

```
C c = new C();
```

```
Thread 1  
c.inc();
```

```
Thread 2  
c.inc();
```

- ❑ A program can often use `this` as the object to lock
- ❑ Does the program above have a data race?
 - No, both threads acquire locks on the same object before they access shared data

Synchronization on this

```
class C {  
    static int cnt;  
    void inc() {  
        synchronized (this) {  
            cnt++;  
        } // end of sync  
    } // end of inc  
  
    void dec() {  
        synchronized (this) {  
            cnt--;  
        } // end of sync  
    } // end of dec  
}
```

```
C c = new C();
```

```
Thread 1  
c.inc();
```

```
Thread 2  
c.dec();
```

- ❑ Does the program above have a data race?
 - No, both threads acquire locks on the same object before they access shared data

Example

□ See

- ShareWelcome/Server.java
- ShareWelcome/ServiceThread.java

Discussion

- ❑ You would not need the lock for `accept` if Java were to label the call as thread safe (synchronized)
- ❑ One reason Java does not specify `accept` as thread safe is that one could register your own socket implementation with [`ServerSocket.setSocketFactory`](#)
- ❑ Always consider thread safety in your design
 - If a resource is shared through concurrent read/write, write/write), consider thread-safe issues.

Why not Synchronization

- ❑ Synchronized method invocations generally are going to be slower than non-synchronized method invocations
- ❑ Synchronization gives rise to the possibility of deadlock, a severe performance problem in which your program appears to hang

Synchronization Overhead

- Try SyncOverhead.java

Synchronization Overhead

□ Try SyncOverhead.java

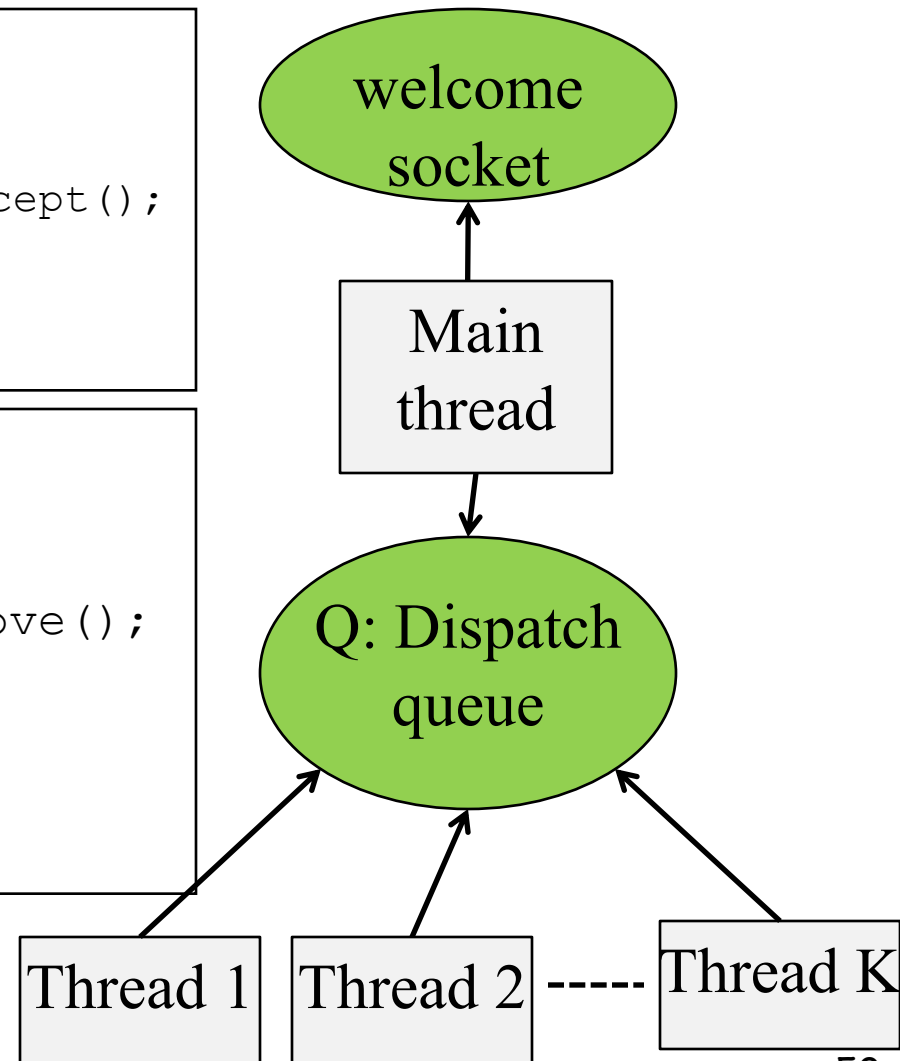
Method	Time (ms; 5,000,000 exec)
no sync	8 ms
synchronized method	18 ms
synchronized on this	18 ms
lock	89 ms
lock and finally	88 ms

Design 2: Producer/Consumer

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```


```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```

How to turn it into
working code?



Main

```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            Q.add(con);  
        } // end of while  
    }  
}
```



```
main {  
    void run {  
        while (true) {  
            Socket con = welcomeSocket.accept();  
            synchronized(Q) {  
                Q.add(con);  
            }  
        } // end of while  
    }  
}
```

Worker

```
WorkerThread {  
    void run {  
        while (true) {  
            Socket myConnSock = Q.remove();  
            // process myConnSock  
            myConnSock.close();  
        } // end of while  
    }  
}
```



```
while (true) {  
    // get next request  
    Socket myConn = null;  
    while (myConn==null) {  
        synchronize(Q) {  
            if (!Q.isEmpty())  
                myConn = (Socket) Q.remove();  
        }  
    } // end of while  
    // process myConn  
}
```

Example

□ try

- ShareQ/Server.java
- ShareQ/ServiceThread.java

Problem of ShareQ Design

- ❑ Worker thread continually spins (**busy wait**) until a condition holds

```
while (true) { // spin
    lock;
    if (Q.condition) // {
        // do something
    } else {
        // do nothing
    }
    unlock
} //end while
```

- ❑ Can lead to high utilization and slow response time
- ❑ Q: Does the shared welcomeSock have busy-wait?