

---

# Network Applications: Load Balancing among Homogeneous Servers; Application Overlays (P2P)

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

11/4/2021

# Outline

---

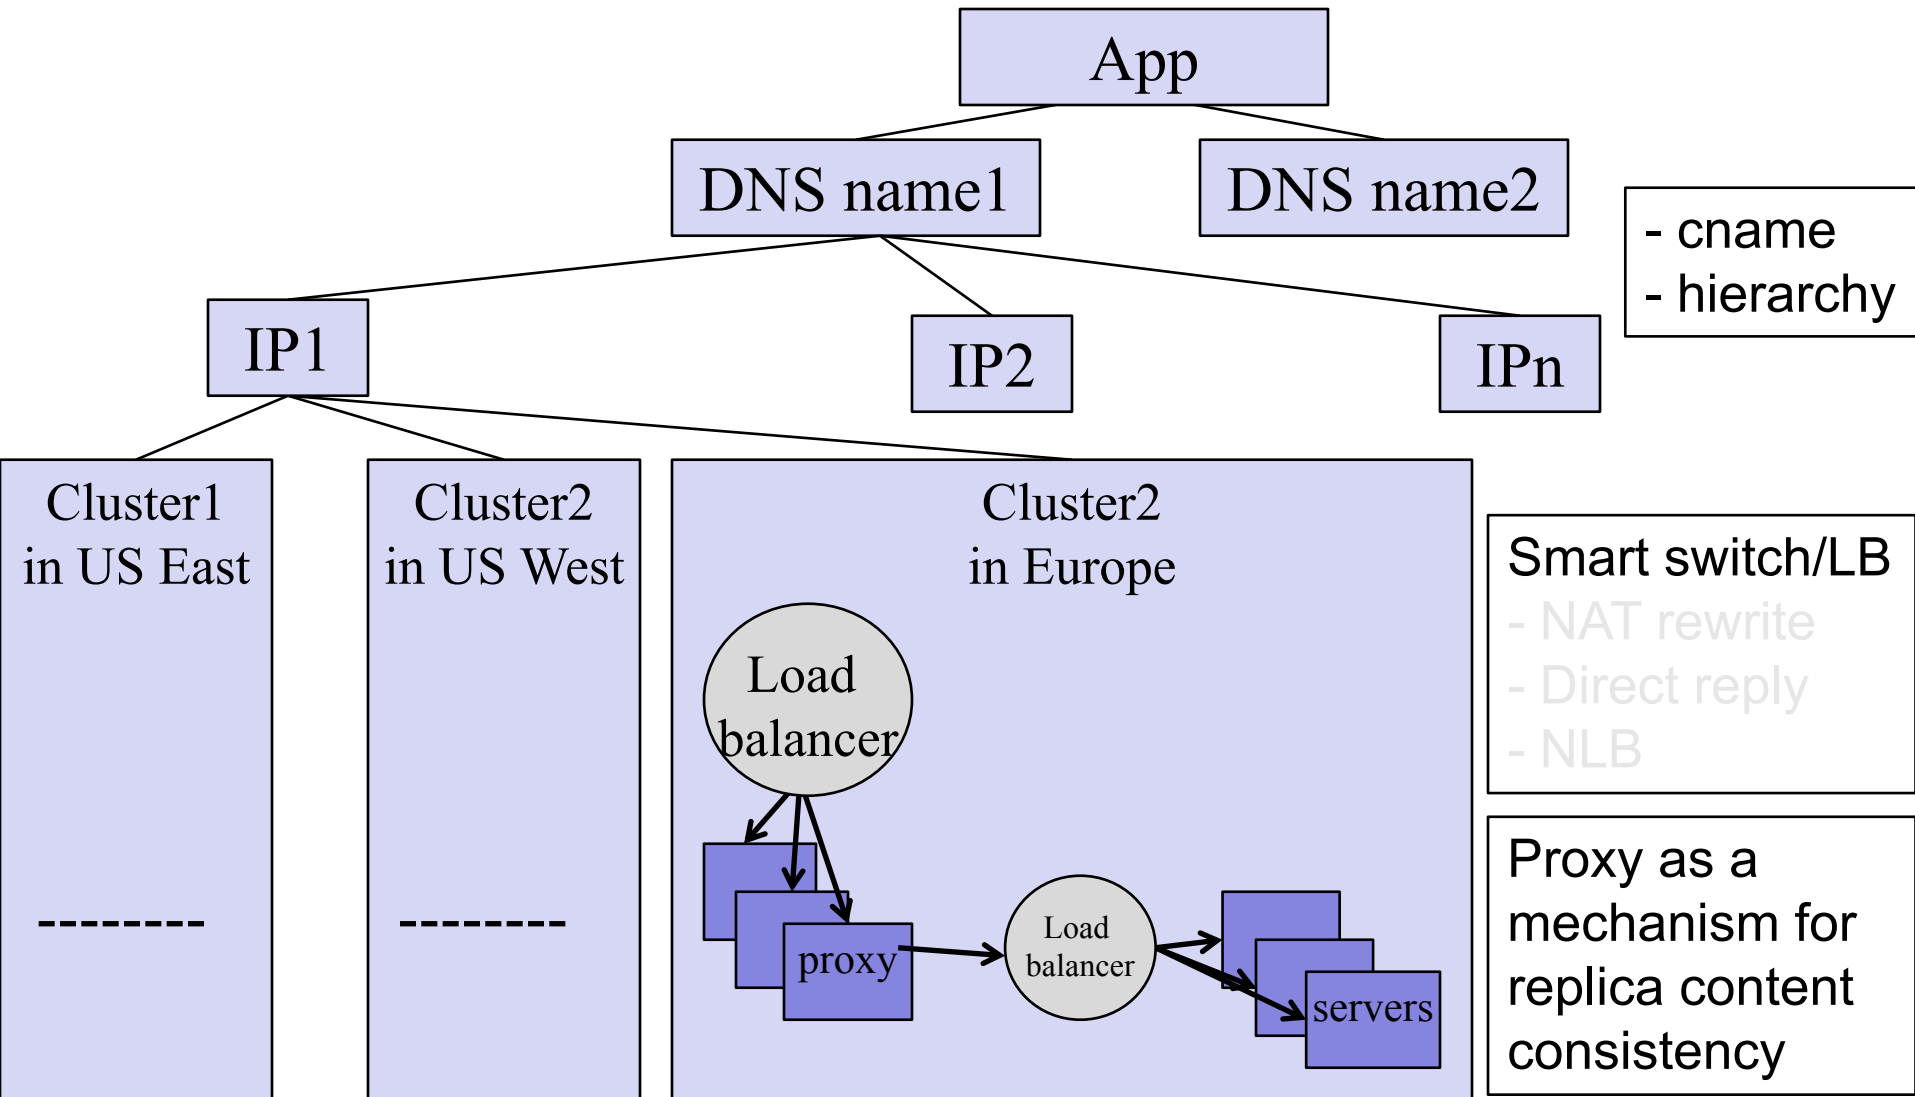
- ❑ Admin and recap
- ❑ Multi-servers
  - Basic issues
  - Load direction
    - DNS (IP level)
    - Load balancer/smart switch (sub-IP level)
- ❑ Application overlays (distributed network applications) to
  - scale bandwidth/resource (BitTorrent)

# Admin

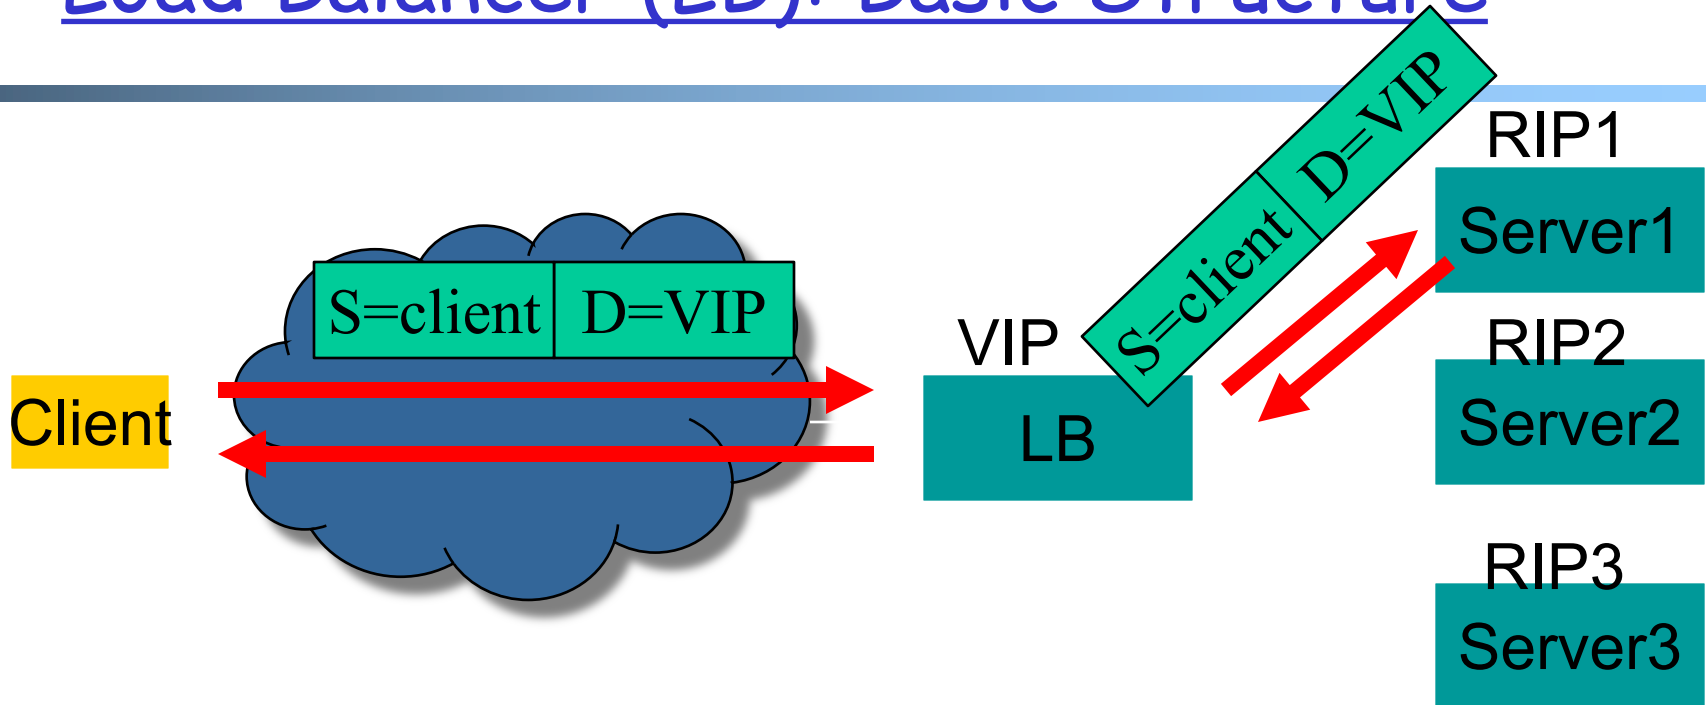
---

- ❑ Assignment three status and questions?
- ❑ Midterm exam 1: Thursday next week (i.e., Nov. 11)
  - cover from lecture 1 to lecture 16 (the one on Nov. 9)

# Recap: Direction Mechanisms



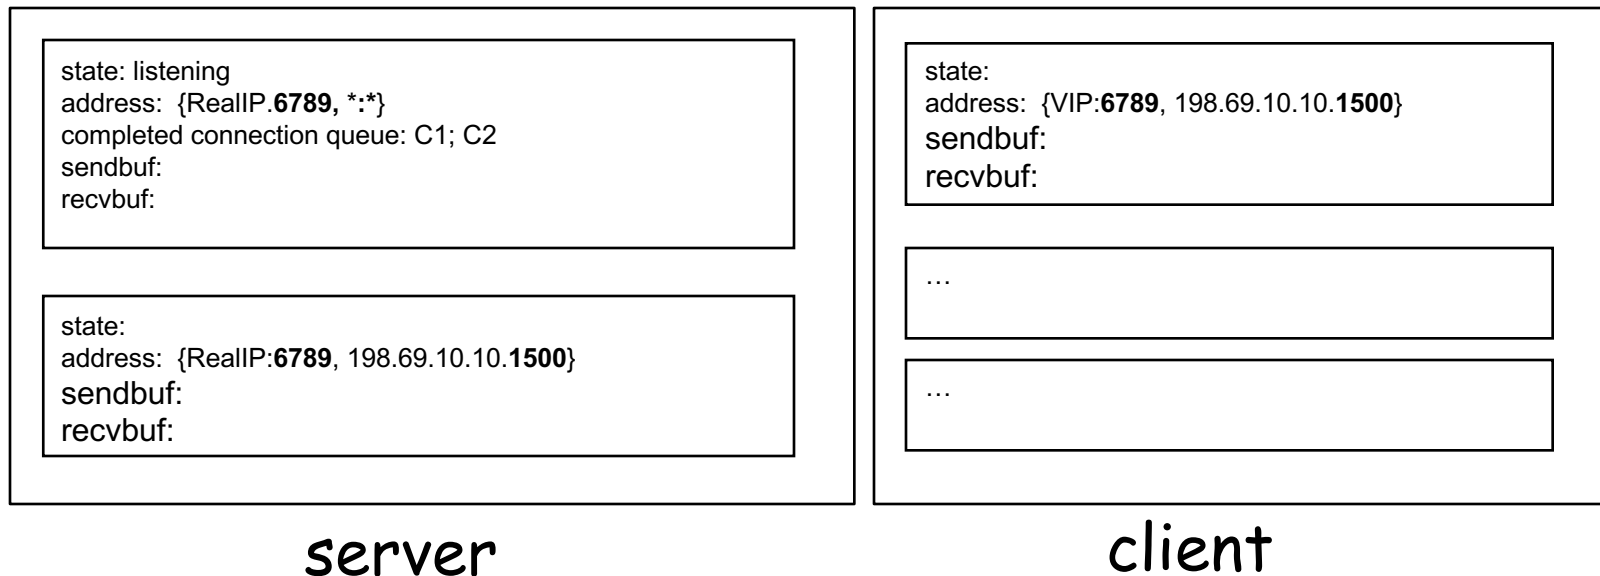
# Load Balancer (LB): Basic Structure



Problem of the basic structure?

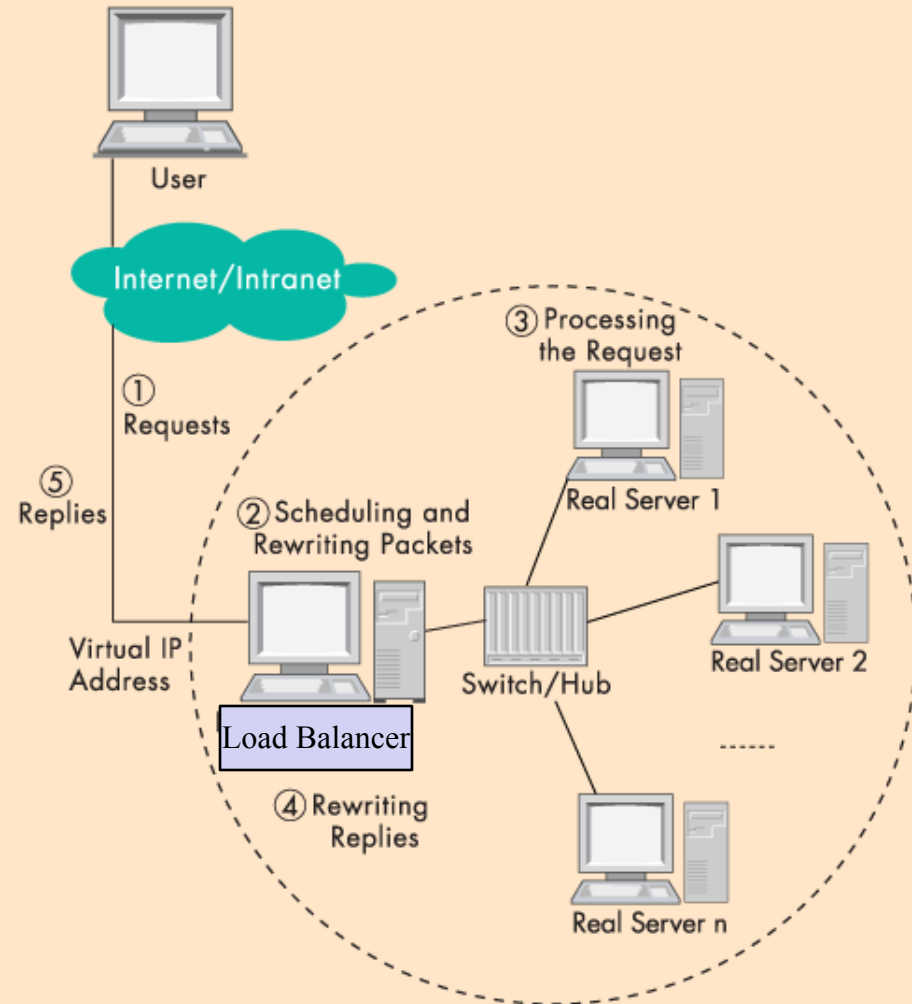
# Problem

- ❑ Client to server packet has VIP as destination address, but real servers use RIPv
  - if LB just forwards the packet from client to a real server, the real server drops the packet
  - reply from real server to client has real server IP as source -> client will drop the packet



# Solution 1: Network Address Translation (NAT)

- ❑ LB does rewriting/translation
- ❑ Thus, the LB is similar to a typical NAT gateway with an additional scheduling function



# Example Virtual Server via NAT


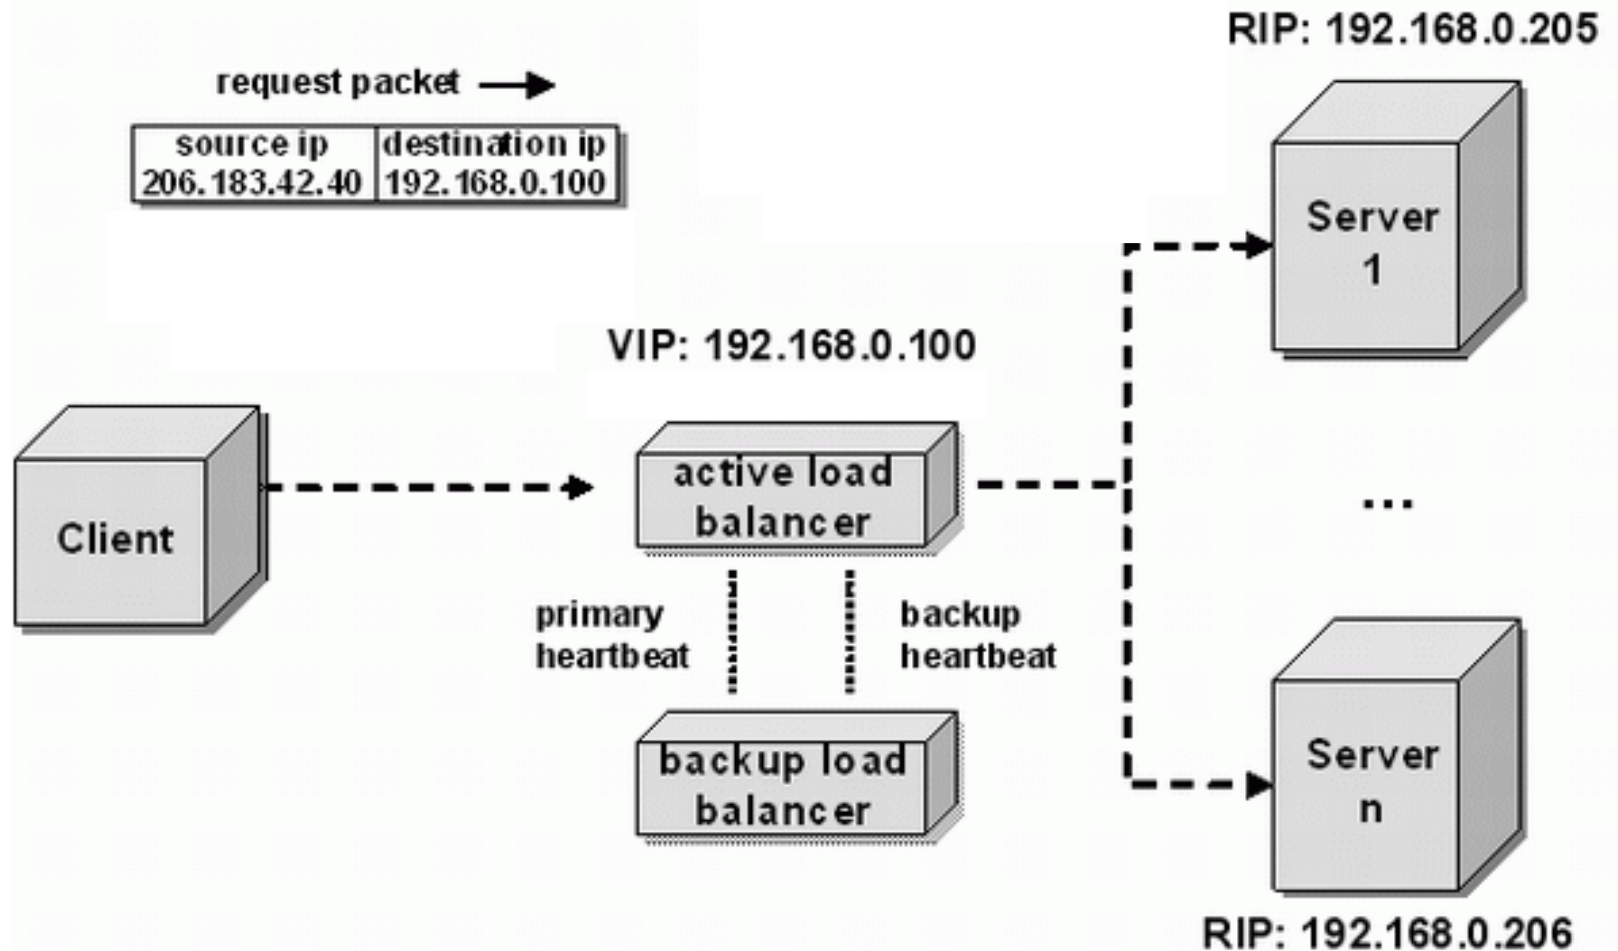


Table 1: an example of virtual server rules

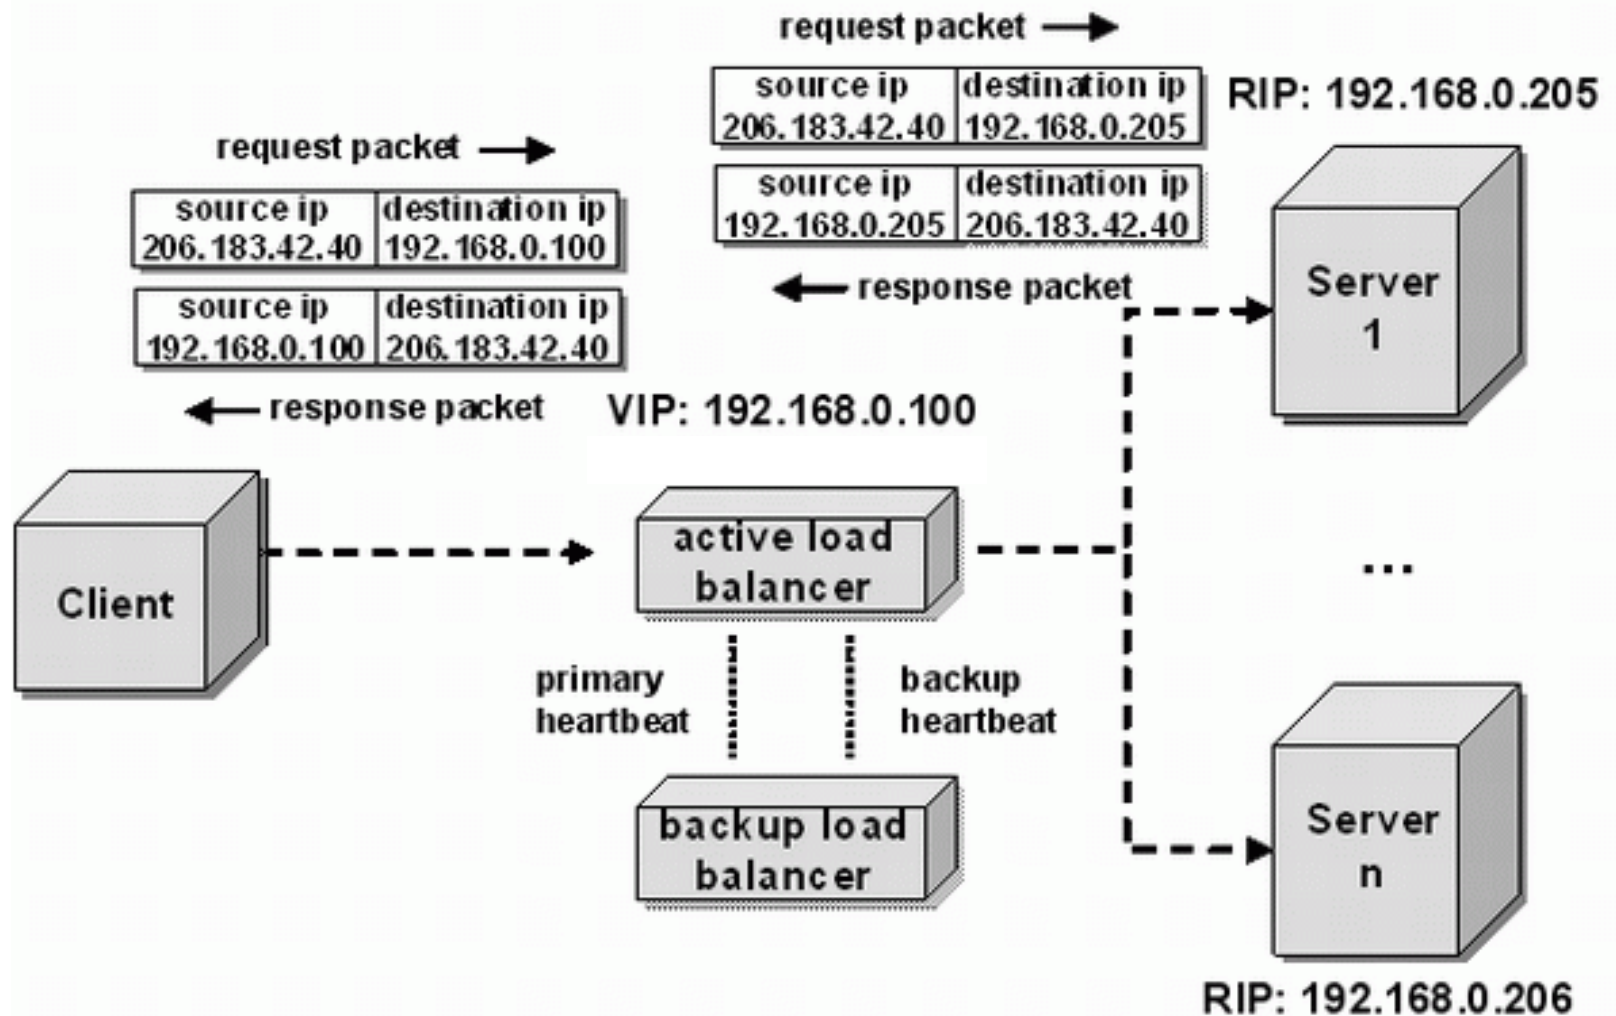
Protocol	Virtual IP Address	Port	Real IP Address	Port	Weight
TCP	202.103.106.5	80	172.16.0.2	80	1
			172.16.0.3	8000	2
TCP	202.103.106.5	21	172.16.0.3	21	1



# LB/NAT Flow



# LB/NAT Flow



# LB/NAT Advantages and Disadvantages

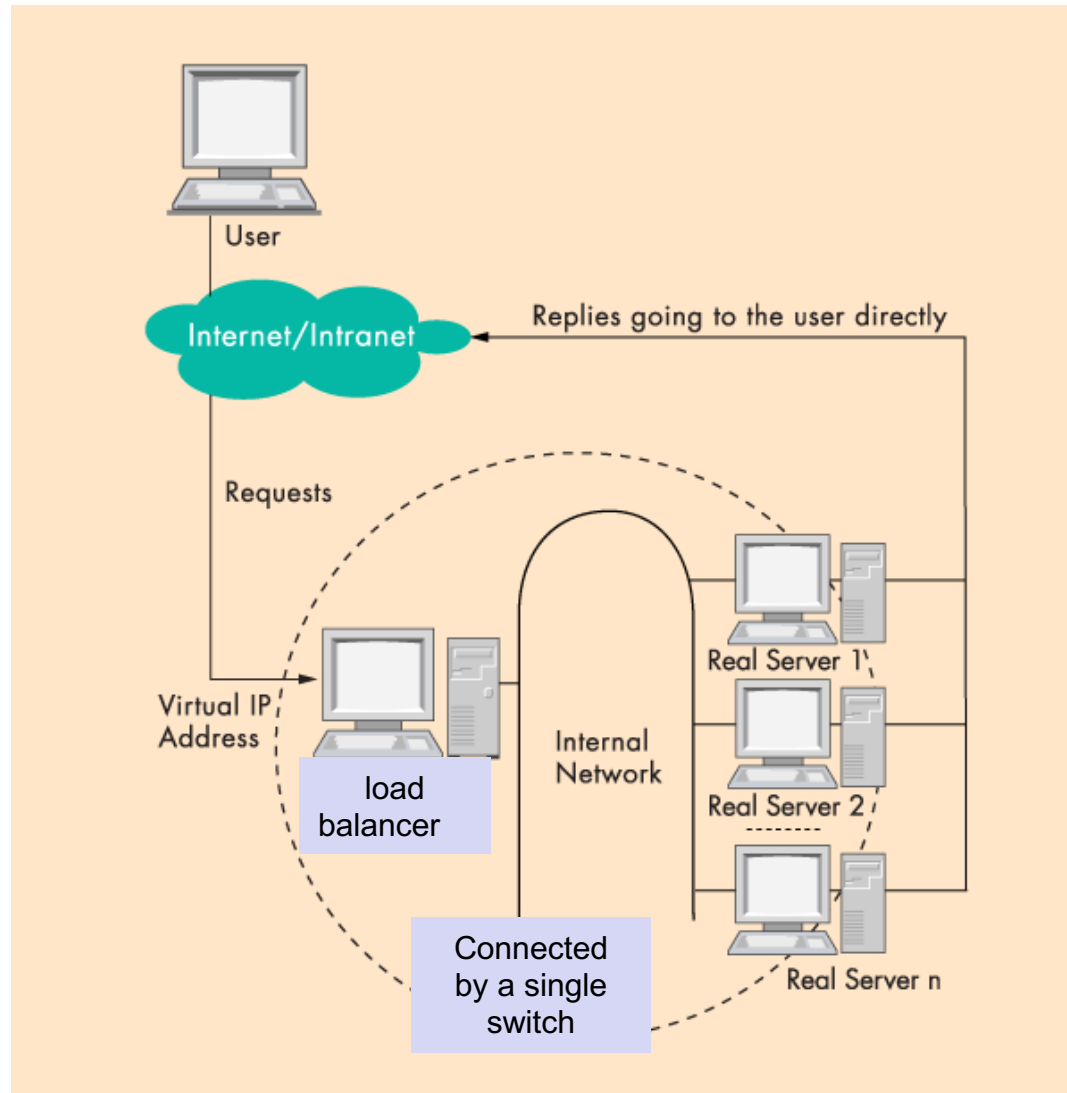
## □ Advantages:

- Only one public IP address is needed for the load balancer; real servers can use private IP addresses
- Real servers need no change and are not aware of load balancing


## □ Problem

- The load balancer must be on the critical path and hence may become the bottleneck due to load to rewrite request and response packets
  - Typically, rewriting responses has more load because there are more response packets

# Goal: LB w/ Direct Reply




# LB with Direct Reply: Implication



- Direct reply → Each real server uses VIP as its IP address
- **Address conflict:** multiple devices w/ the same IP addr

# Why IP Address Matters?




- ❑ Each network interface card listens to an assigned MAC address
- ❑ A router is configured with the range of IP addresses connected to each interface (NIC)
- ❑ To send to a device with a given IP, the router needs to translate IP to MAC (device) address
- ❑ The translation is done by the Address Resolution Protocol (ARP)

# ARP Protocol

- ❑ ARP is “plug-and-play”:
  - nodes create their ARP tables without intervention from net administrator
- ❑ A **broadcast** protocol:
  - Router broadcasts query frame, containing queried IP address
    - all machines on LAN receive ARP query
  - Node with queried IP receives ARP frame, replies its MAC address


# ARP in Action



- Router broadcasts ARP broadcast query: who has VIP?
- ARP reply from LB: I have VIP; my MAC is  $MAC_{LB}$
- Data packet from R to LB: destination MAC =  $MAC_{LB}$



# LB/DR Problem



ARP and race condition:

- When router R gets a packet with dest. address VIP, it broadcasts an Address Resolution Protocol (ARP) request: who has VIP?
- One of the real servers may reply before load balancer

Solution: configure real servers to not respond to ARP request

# LB via Direct Routing

- ❑ The virtual IP address is shared by real servers and the load balancer.
- ❑ Each real server has a **non-ARPing**, loopback alias interface configured with the virtual IP address, and the load balancer has an interface configured with the virtual IP address to accept incoming packets.
- ❑ The workflow of LB/DR is similar to that of LB/NAT:
  - the load balancer directly routes a packet to the selected server
    - the load balancer simply changes the MAC address of the data frame to that of the server and retransmits it on the LAN (how to know the real server's MAC?)
  - When the server receives the forwarded packet, the server determines that the packet is for the address on its loopback alias interface, processes the request, and finally returns the result directly to the user

# LB/DR Advantages and Disadvantages

---

## ❑ Advantages:

- Real servers send response packets to clients directly, avoiding LB as bottleneck


## ❑ Disadvantages:

- Servers must have non-arp alias interface
- The load balancer and server must have one of their interfaces in the same LAN segment
- Considered by some as a hack, not a clean architecture

# Example Implementation of LB

- ❑ An example open source implementation is Linux virtual server ([linux-vs.org](http://linux-vs.org))
  - Used by
    - [www.linux.com](http://www.linux.com)
    - [sourceforge.net](http://sourceforge.net)
    - [wikipedia.org](http://wikipedia.org)
  - More details on ARP problem:  
[http://www.austintek.com/LVS/LVS-HOWTO/HOWTO/LVS-HOWTO.arp\\_problem.html](http://www.austintek.com/LVS/LVS-HOWTO/HOWTO/LVS-HOWTO.arp_problem.html)
  - Many commercial LB servers from F5, Cisco, ...
- ❑ More details please read chapter 2 of [Load Balancing Servers, Firewalls, and Caches](#)


# Problem of the Load Balancer Architecture



One major problem is that the LB becomes a single point of failure (SPOF).

# Solutions

- ❑ Redundant load balancers
  - E.g., two load balancers (a good question to think offline)
- ❑ Fully distributed load balancing
  - e.g., Microsoft Network Load Balancing (NLB)



# Microsoft NLB

- ❑ No dedicated load balancer
- ❑ All servers in the cluster receive all packets
- ❑ Key issue: one and only one server processes each packet
  - ❑ All servers within the cluster simultaneously run a mapping algorithm to determine which server should handle the packet. Those servers not required to service the packet simply discard it.
  - ❑ Mapping (ranking) algorithm: computing the “winning” server according to host priorities, multicast or unicast mode, port rules, affinity, load percentage distribution, client IP address, client port number, other internal load information

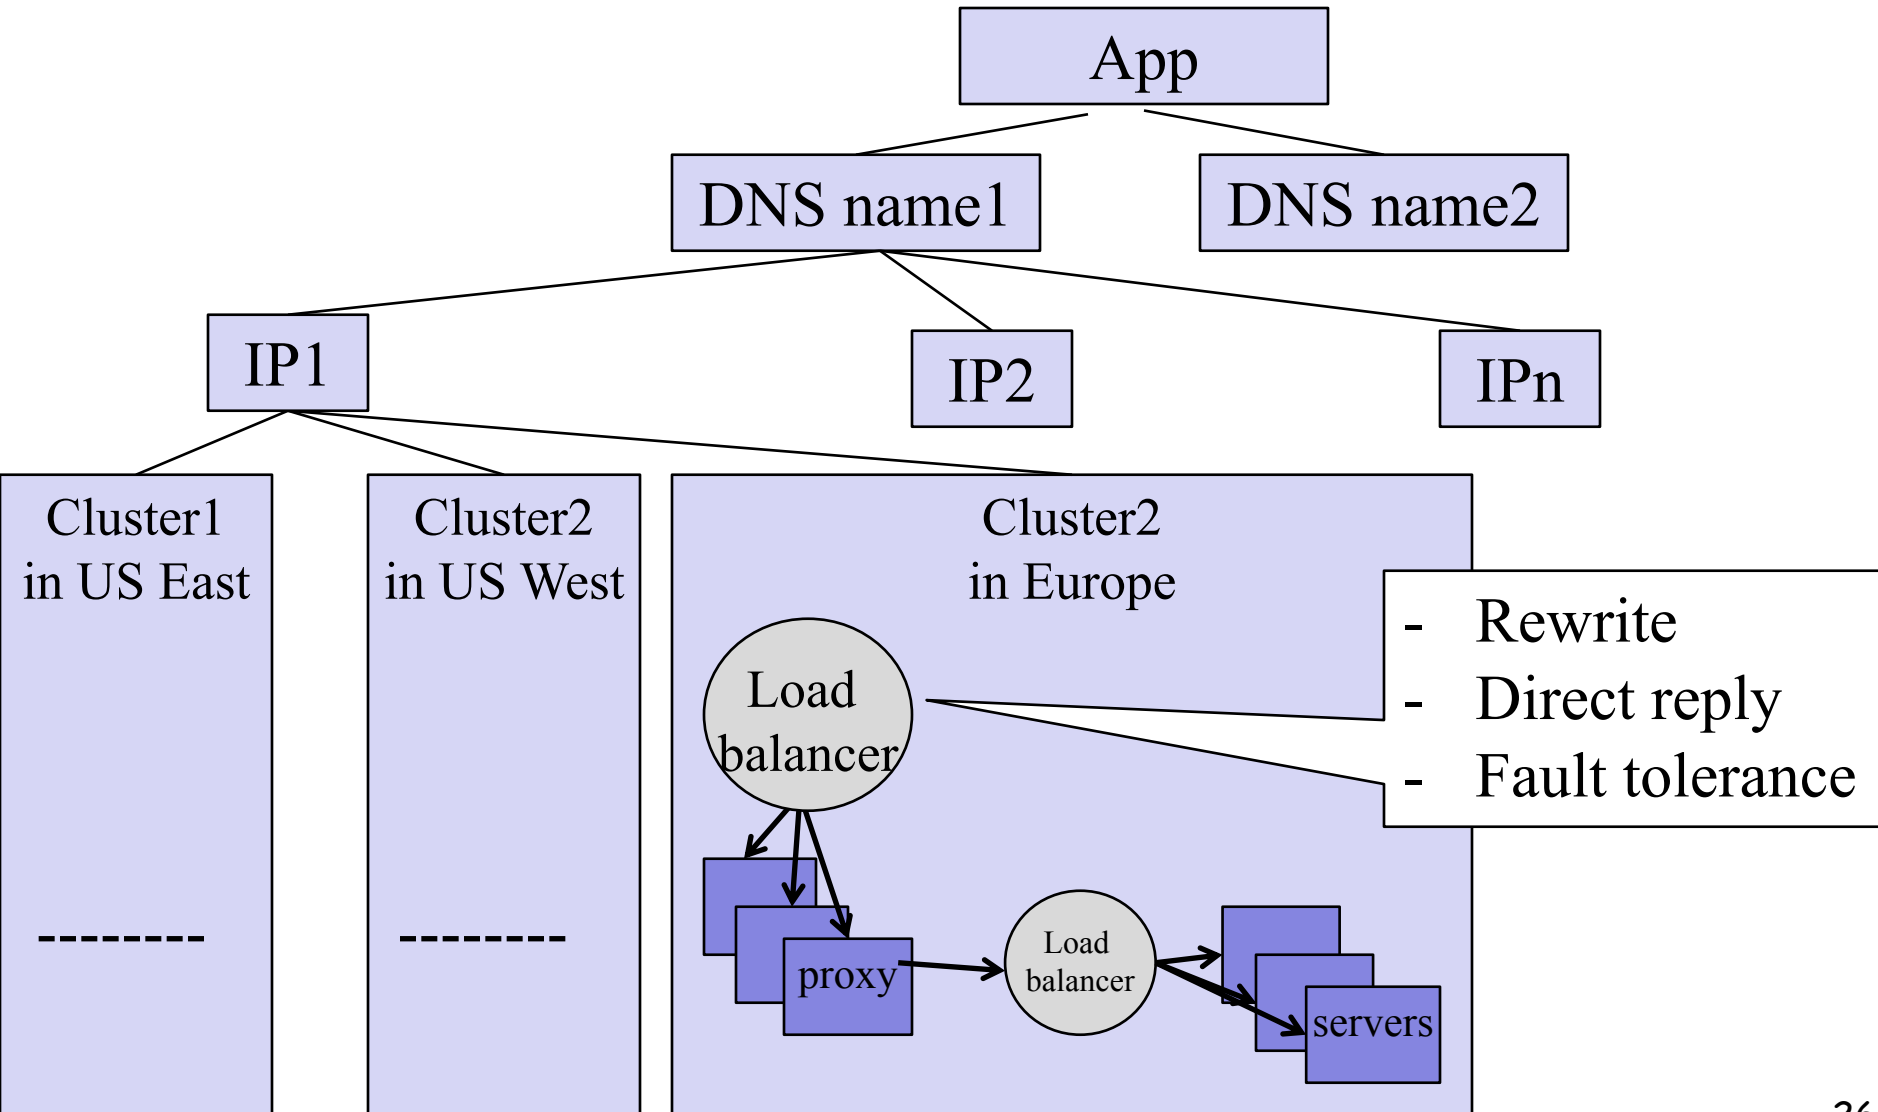
# Discussion

---

- ❑ Compare the design of using Load Balancer vs Microsoft NLB



# Recap: Direction Mechanisms



# Outline

---

- ❑ Admin and recap
- ❑ Single, high-performance network server
- ❑ Multiple servers
  - Overview
  - Basic mechanisms
  - Example: YouTube (offline read)

# You Tube

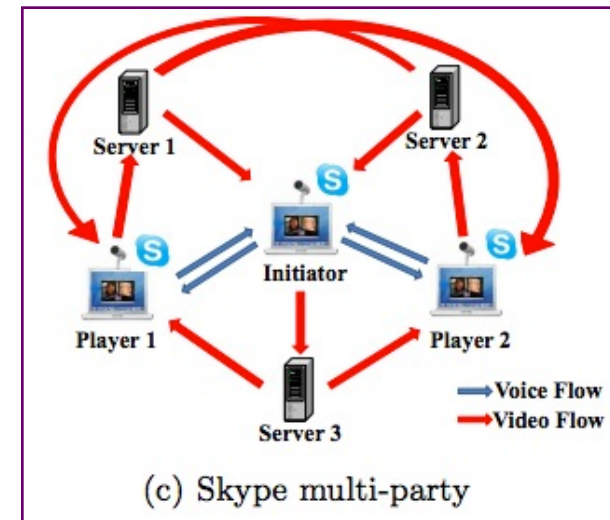

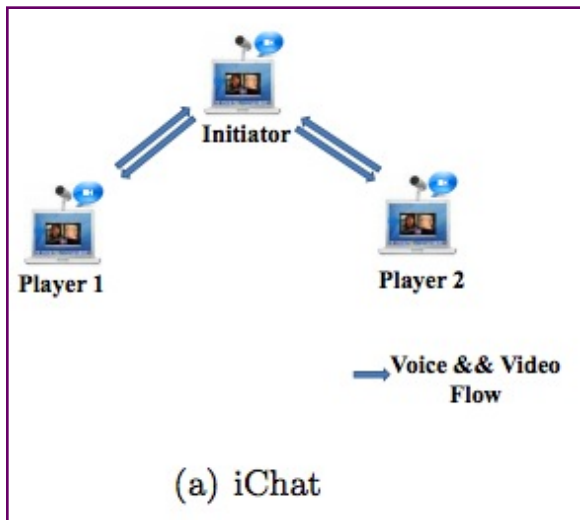
- ❑ 02/2005: Founded by Chad Hurley, Steve Chen and Jawed Karim, who were all early employees of PayPal.
- ❑ 10/2005: First round of funding (\$11.5 M)
- ❑ 03/2006: 30 M video views/day
- ❑ 07/2006: 100 M video views/day
- ❑ 11/2006: acquired by Google
- ❑ 10/2009: Chad Hurley announced in a blog that YouTube serving well over 1 B video views/day (avg = 11,574 video views /sec )

# Pre-Google Team Size


---

- ❑ 2 Sysadmins
- ❑ 2 Scalability software architects
- ❑ 2 feature developers
- ❑ 2 network engineers
- ❑ 1 DBA
- ❑ 0 chefs

# Example Server Systems



# Scalability of Server-Only Approaches




# Outline

---

- ❑ Admin and recap
- ❑ Multiple servers
- ❑ Application overlays
  - potential

# An Upper Bound on Scalability

- ❑ Idea: use resources from both clients and the server
- ❑ Assume
  - need to achieve same rate to all clients
  - only uplinks can be bottlenecks
- ❑ What is an upper bound on scalability?






# The Scalability Problem

- Maximum throughput

$$R = \min\{C_0, (C_0 + \sum C_i)/n\}$$

- The bound is theoretically approachable



# Theoretical Capacity: upload is bottleneck

$$R = \min\{C_0, (C_0 + \sum C_i)/n\}$$

□ Assume  $C_0 > (C_0 + \sum C_i)/n$

□ Tree i:

server  $\rightarrow$  client i:  $C_i/(n-1)$



client i  $\rightarrow$  other  $n-1$  clients

□ Tree 0:



server has remaining

$$C_m = C_0 - (C_1 + C_2 + \dots + C_n)/(n-1)$$

send to client i:  $C_m/n$



# Why not Building the Trees?



- ❑ Clients come and go (churns): maintaining the trees is too expensive
- ❑ Each client needs  $N$  connections

# Server+Host (P2P) Content Distribution: Key Design Issues

## ❑ Robustness


- Resistant to churns and failures

## ❑ Efficiency

- A client has content that others need; otherwise, its upload capacity may not be utilized

## ❑ Incentive: clients are willing to upload

- Some real systems nearly 50% of all responses are returned by the top 1% of sharing hosts




# Discussion: How to handle the issues?

□ Robustness

□ Efficiency

□ Incentive




# Example: BitTorrent

---

- ❑ A P2P file sharing protocol
- ❑ Created by Bram Cohen in 2004
  - Spec at bep\_0003:  
[http://www.bittorrent.org/beps/bep\\_0003.html](http://www.bittorrent.org/beps/bep_0003.html)

# BitTorrent: Lookup



# Metadata (.torrent) File Structure

---

- ❑ Meta info contains information necessary to contact the tracker and describes the files in the torrent
  - URL of tracker
  - file name
  - file length
  - piece length (typically 256KB)
  - SHA-1 hashes of pieces for verification
  - also creation date, comment, creator, ...



# Tracker Protocol

- ❑ Communicates with clients via HTTP/HTTPS
- ❑ Client GET request
  - info\_hash: uniquely identifies the file
  - peer\_id: chosen by and uniquely identifies the client
  - client IP and port
  - numwant: how many peers to return (defaults to 50)
  - stats: e.g., bytes uploaded, downloaded
- ❑ Tracker GET response
  - interval: how often to contact the tracker
  - list of peers, containing peer id, IP and port
  - stats

# Tracker Protocol



“register”

list of peers

ID1	169.237.234.1:6881
ID2	190.50.34.6:5692
ID3	34.275.00.142:4545
...	
ID50	23.145.14.145:6881




# Robustness and efficiency:

## Piece-based Swarming

- Divide a large file into small blocks and request block-size content from different peers (why?)


Block: unit of download



- If do not finish downloading a block from one peer within timeout (say due to churns), switch to requesting the block from another peer

# Detail: Peer Protocol


(Over TCP)



- ❑ Peers exchange bitmap representing content availability
  - `bitfield` msg during initial connection
  - `have` msg to notify updates to bitmap
  - to reduce bitmap size, aggregate multiple blocks as a piece

# Peer Request

- ❑ If peer A has a piece that peer B needs, peer B sends `interested` to A
- ❑ `unchoke`: indicate that A allows B to request
- ❑ `request`: B requests a specific block from A
- ❑ `piece`: specific data




# Key Design Points

## □ request:

- which data blocks to request?

## □ unchoke:

- which peers to serve?



# Request: Block Availability

---

- Request (local) **rarest first**
  - achieves the fastest replication of rare pieces
  - obtain something of value

# Block Availability: Revisions

- ❑ When downloading starts (first 4 pieces): choose at random and request them from the peers
  - get pieces as quickly as possible
  - obtain something to offer to others
  
- ❑ Endgame mode
  - defense against the “last-block problem”: cannot finish because missing a few last pieces
  - send requests for missing pieces to all peers in our peer list
  - send `cancel` messages upon receipt of a piece




# BitTorrent: Unchoke

- Periodically (typically every 10 seconds) calculate data-receiving rates from all peers

- Upload to (*unchoke*) the fastest

- constant number (4) of unchoking slots
- partition upload bw equally among unchoked

commonly referred to as “**tit-for-tat**” strategy



# Optimistic Unchoking

---

- ❑ Periodically select a peer at random and upload to it
  - typically every 3 unchoking rounds (30 seconds)
- ❑ Multi-purpose mechanism
  - allow bootstrapping of new clients
  - continuously look for the fastest peers (exploitation vs exploration)

# BitTorrent Fluid Analysis

- Normalize file size to 1
- $x(t)$ : number of downloaders (also known as leechers) who do not have all pieces at time  $t$ .
- $y(t)$ : number of seeds in the system at time  $t$ .
- $\lambda$ : the arrival rate of new requests.
- $\mu$ : the uploading bandwidth of a given peer.
- $c$ : the downloading bandwidth of a given peer, assume  $c \geq \mu$ .
- $\theta$ : the rate at which downloaders abort download.
- $\gamma$ : the rate at which seeds leave the system.
- $\eta$ : indicates the effectiveness of downloader sharing,  $\eta$  takes values in  $[0, 1]$ .

# System Evolution

$$\begin{aligned}\frac{dx}{dt} &= \lambda - \theta x(t) - \min\{cx(t), \mu(\eta x(t) + y(t))\}, \\ \frac{dy}{dt} &= \min\{cx(t), \mu(\eta x(t) + y(t))\} - \gamma y(t),\end{aligned}$$

Solving steady state:  $\frac{dx(t)}{dt} = \frac{dy(t)}{dt} = 0$


Define  $\frac{1}{\beta} = \max\left\{\frac{1}{c}, \frac{1}{\eta}\left(\frac{1}{\mu} - \frac{1}{\gamma}\right)\right\}$

$$\begin{aligned}\bar{x} &= \frac{\lambda}{\beta(1 + \frac{\theta}{\beta})} \\ \bar{y} &= \frac{\lambda}{\gamma(1 + \frac{\theta}{\beta})}.\end{aligned}$$

# System State

$$\bar{x} = \frac{\lambda}{\beta(1 + \frac{\theta}{\beta})}$$
$$\bar{y} = \frac{\lambda}{\gamma(1 + \frac{\theta}{\beta})}$$

Q: How long does each downloader stay as a downloader?



$$T = \frac{1}{\theta + \beta}$$

$$\frac{1}{\beta} = \max\left\{\frac{1}{c}, \frac{1}{\eta}\left(\frac{1}{\mu} - \frac{1}{\gamma}\right)\right\}$$

Key take-away: not scaling inverse with system size (x)

- New requests comes, new bandwidth also comes