
Network Transport Layer:
Transport Reliability:
Sliding Windows; Connection Management; TCP

Qiao Xiang

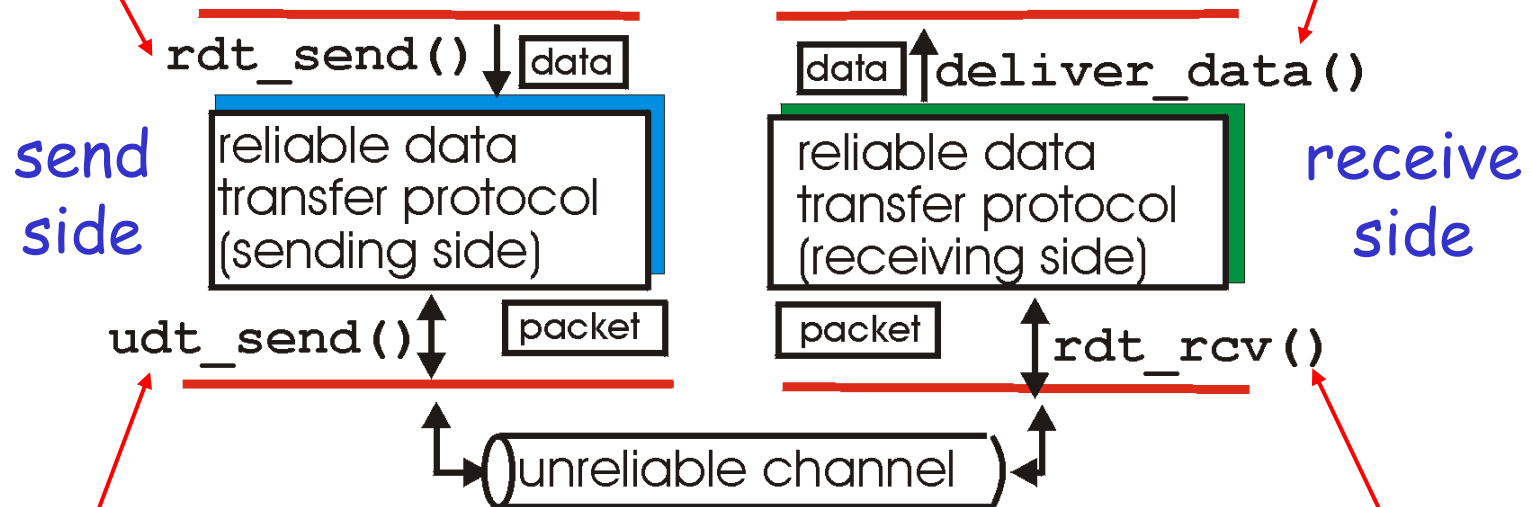
<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

11/11/2021

Recap: Reliable Data Transfer Context

rdt_send() : called from above,
(e.g., by app.)

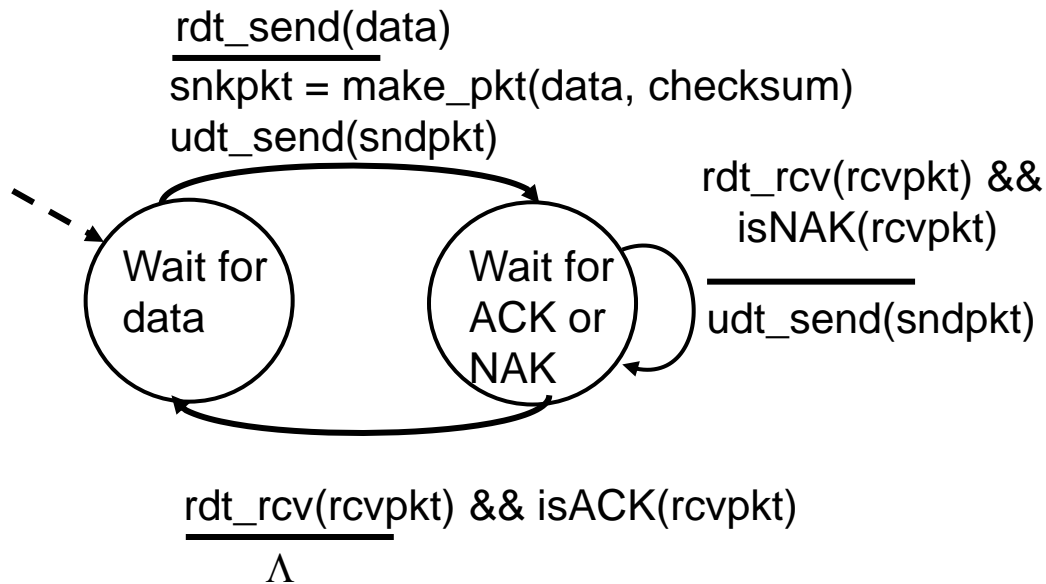
deliver_data() : called by
rdt to deliver data to upper



udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver


rdt_rcv() : called from below;
when packet arrives on rcv-side of
channel

Recap: rdt2.0: Reliability allowing only Data Msg Corruption

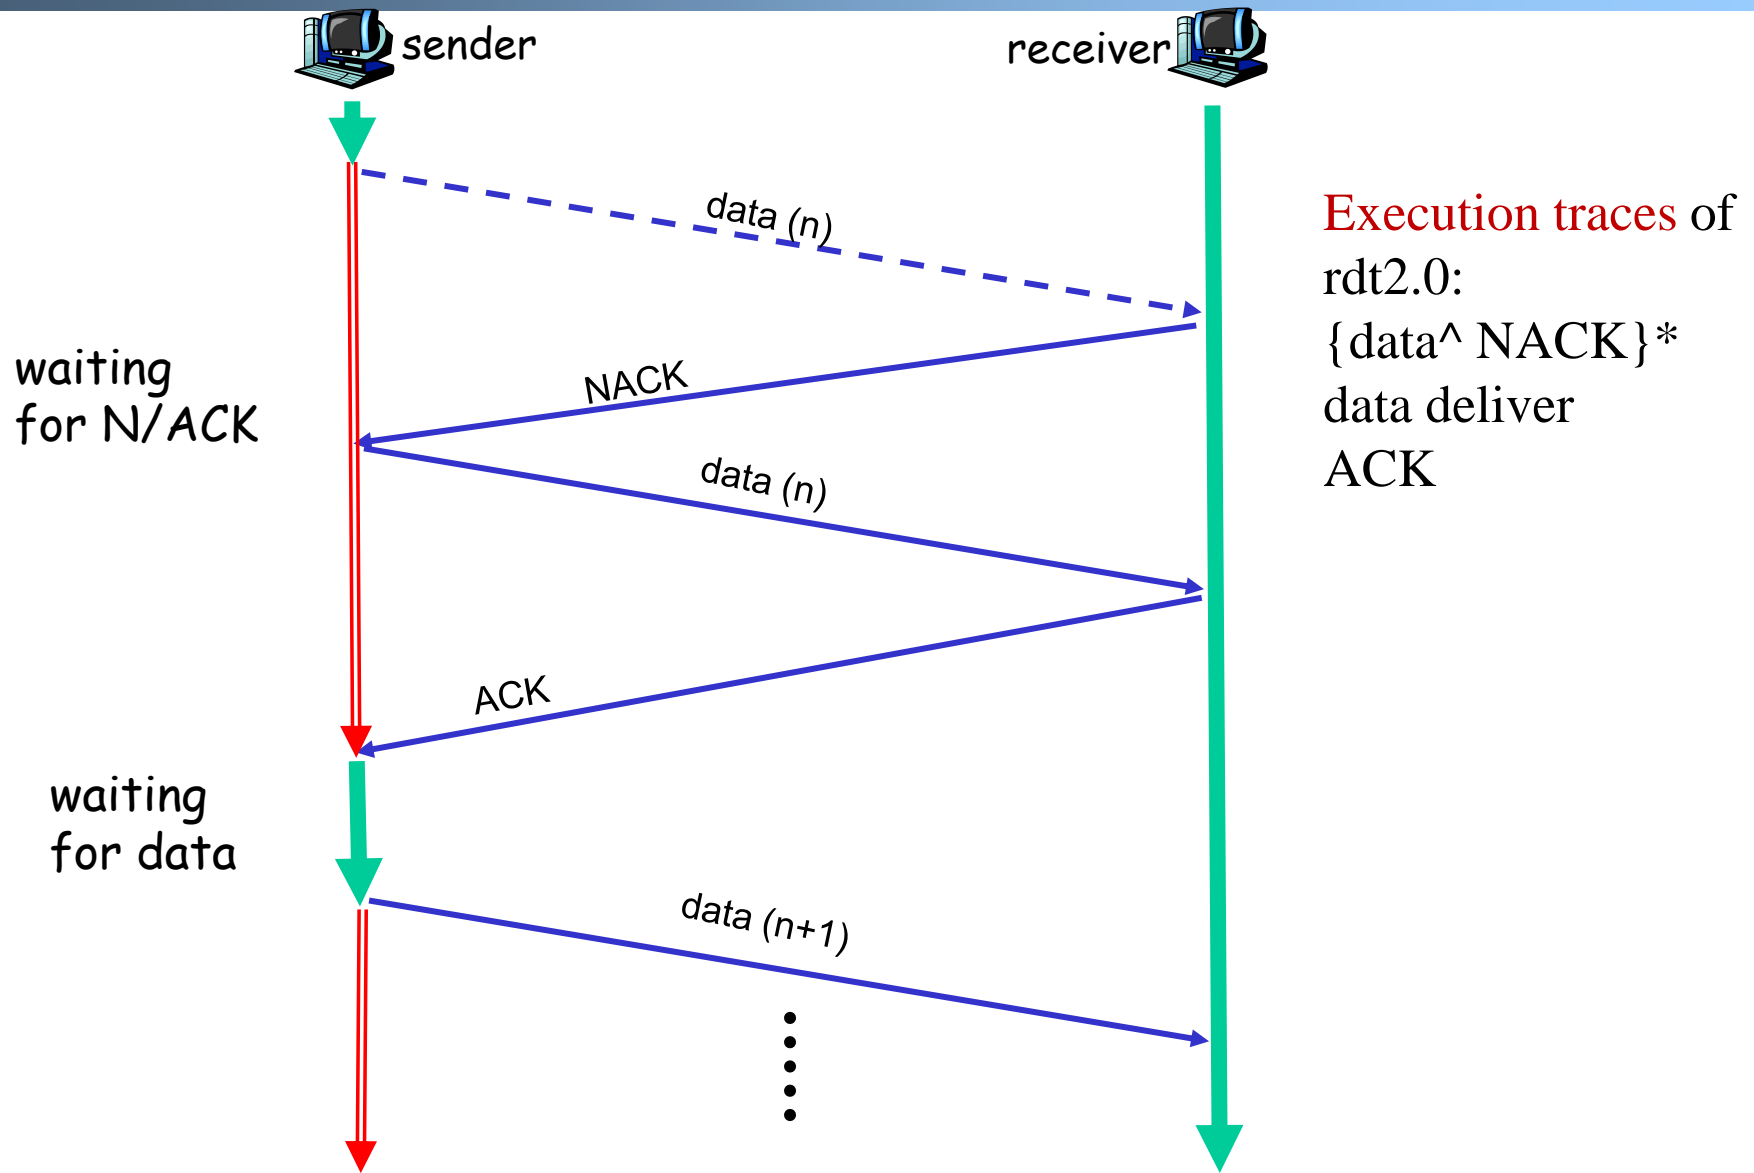


sender

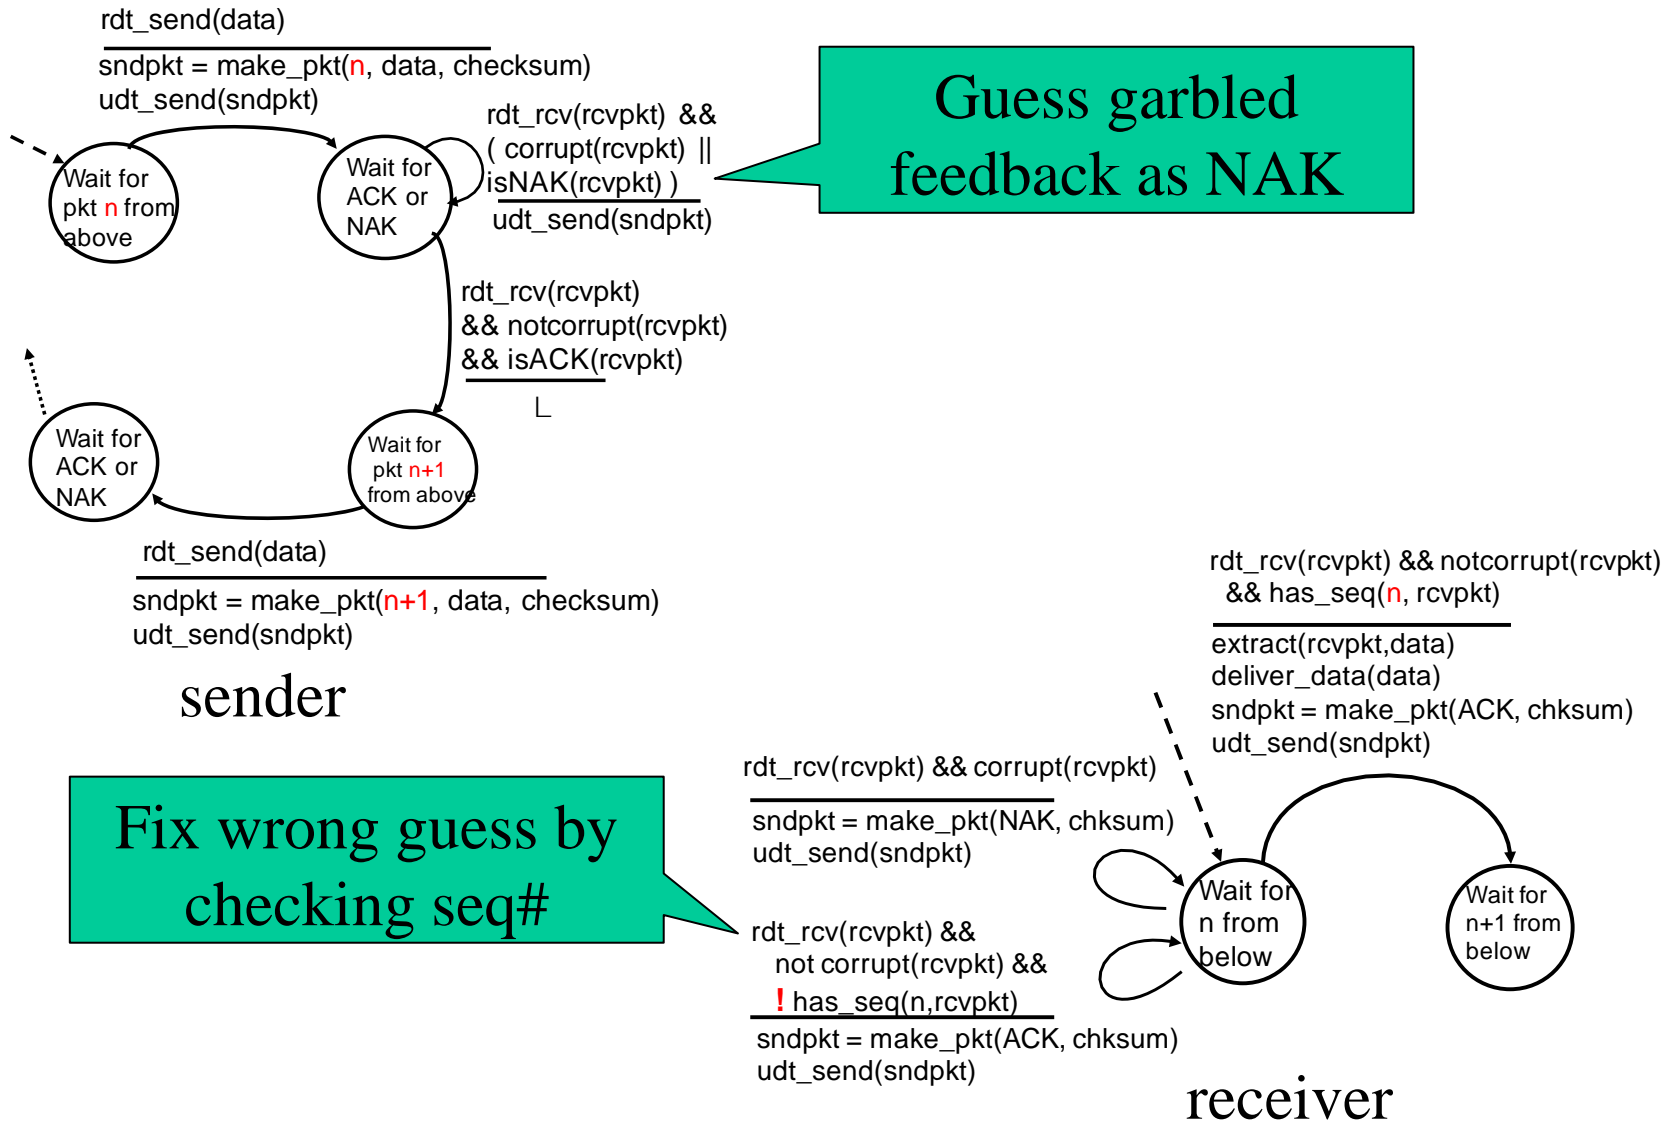
receiver



Recap: Rdt2.0 Analysis




Recap: rdt2.1b: Reliability allowing Data/Control Msg Corruption



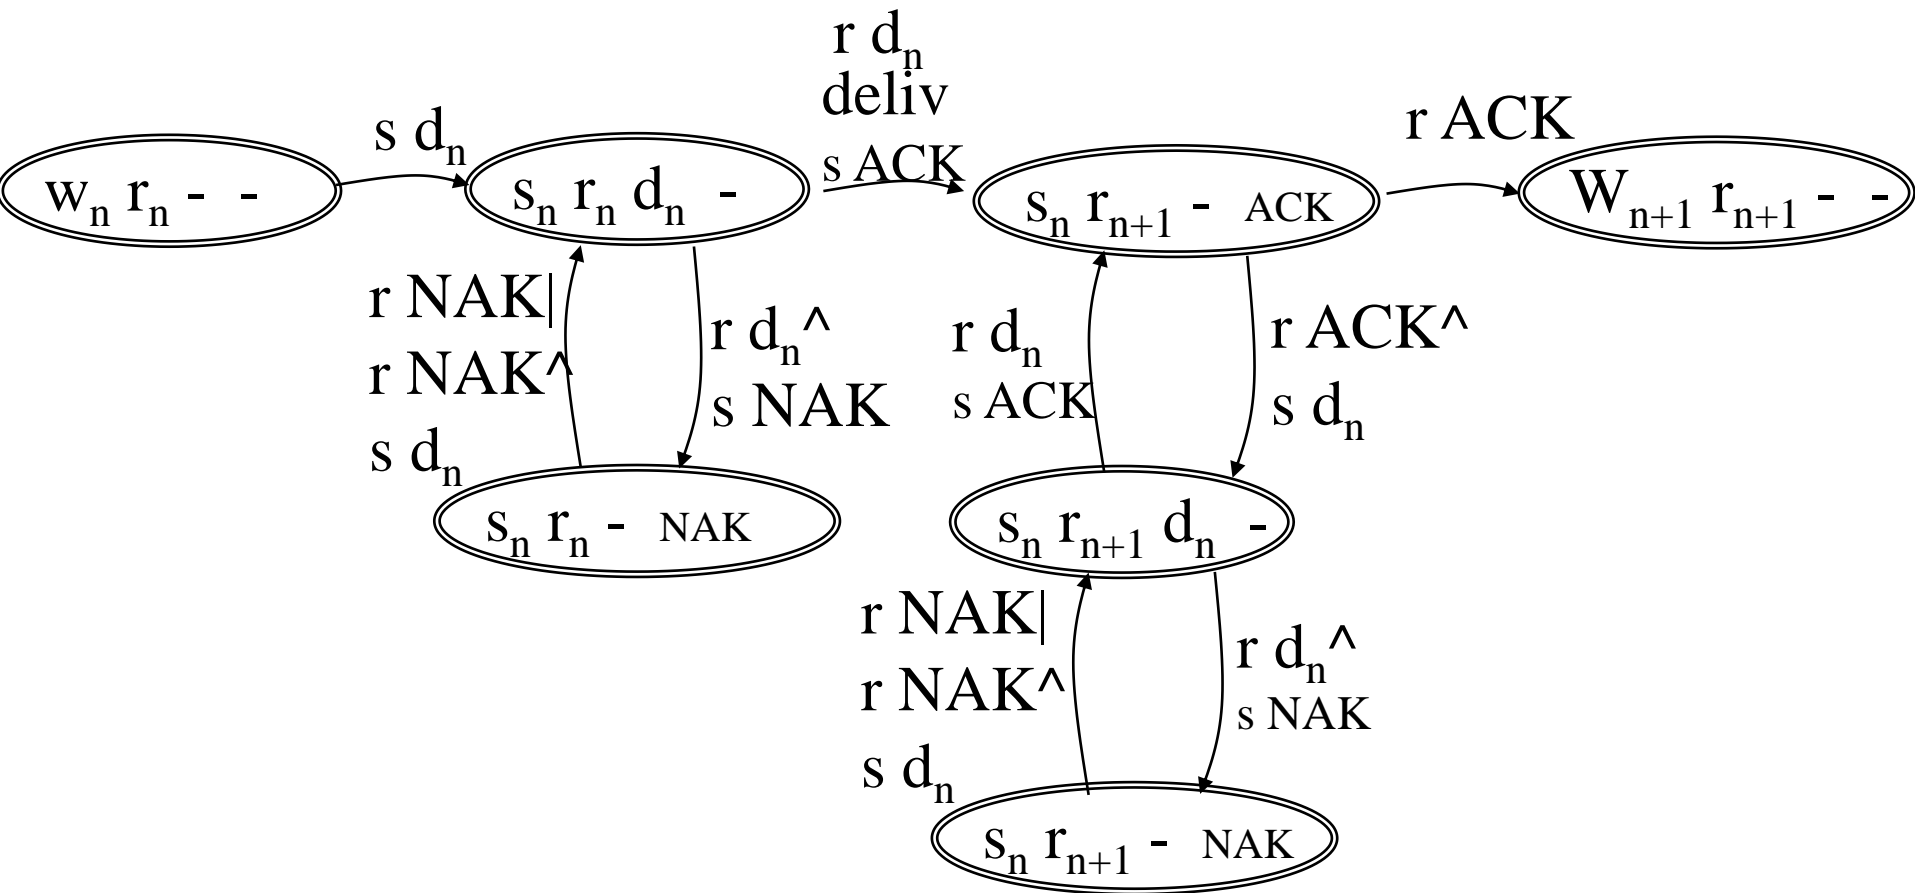
Protocol Analysis using (Generic) Execution Traces Technique

- ❑ Issue: how to systematically enumerate all potential execution traces to understand and verify correctness
- ❑ A systematic approach to enumerating exec. traces is to compute **joint sender/receiver/channels state machine**

sender state: waiting for n

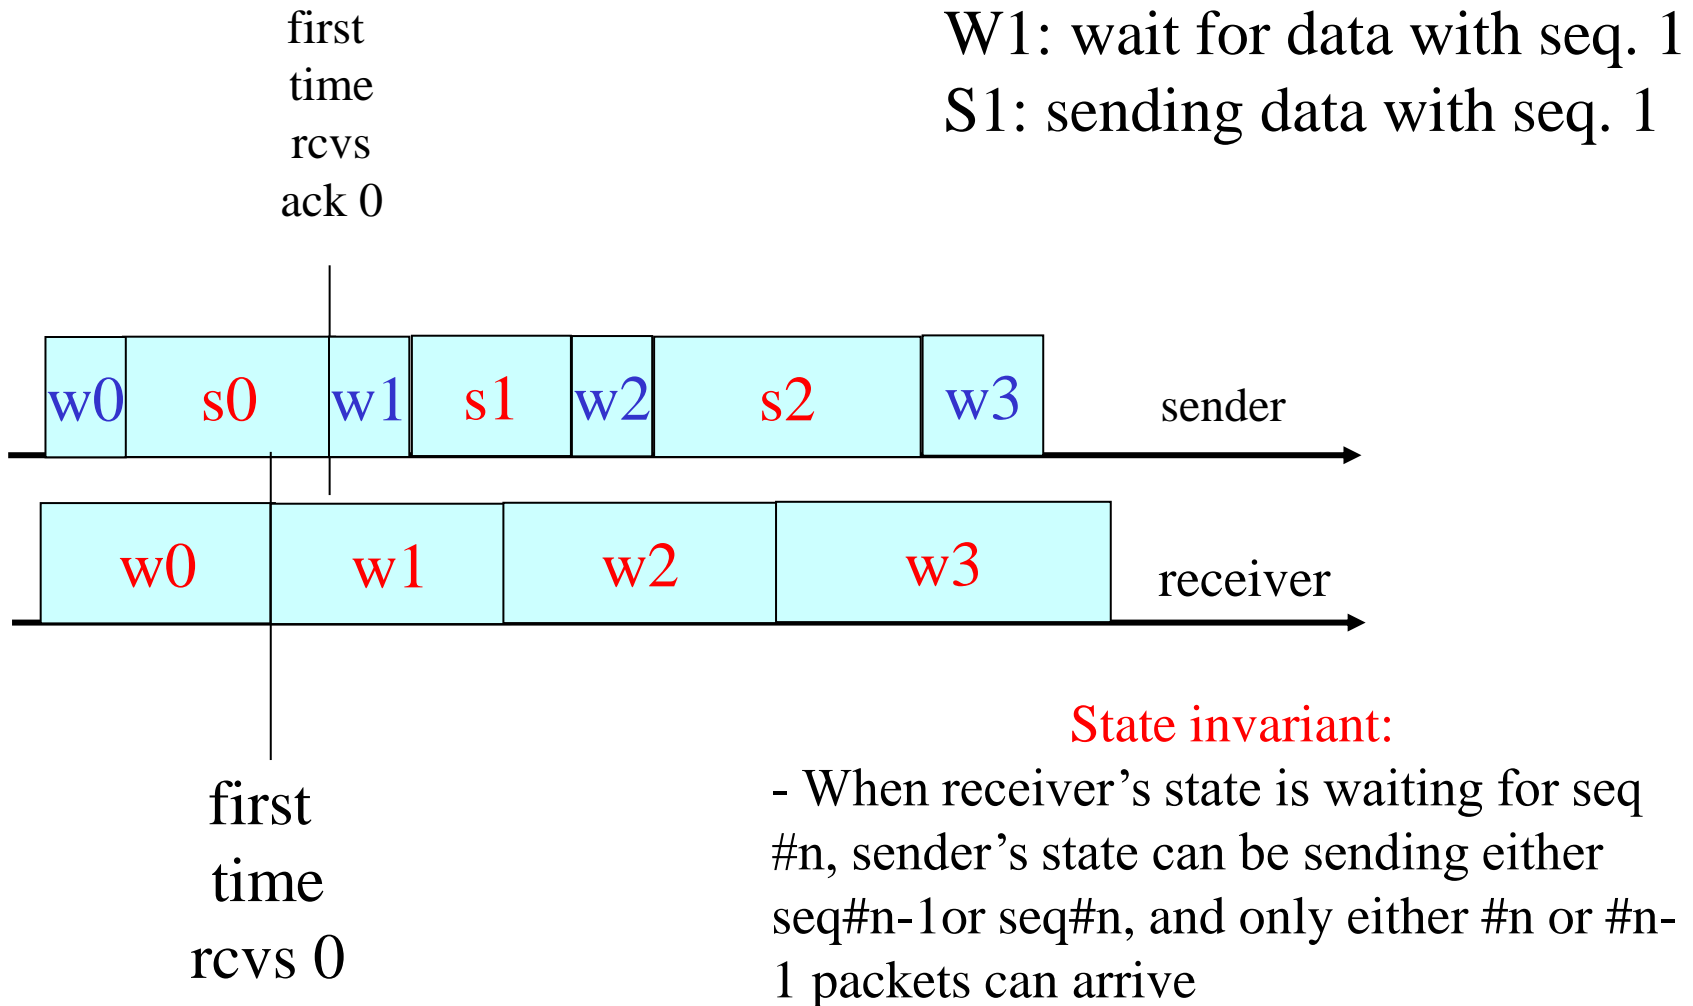


Recap: Protocol Analysis using (Generic) Execution Traces Technique

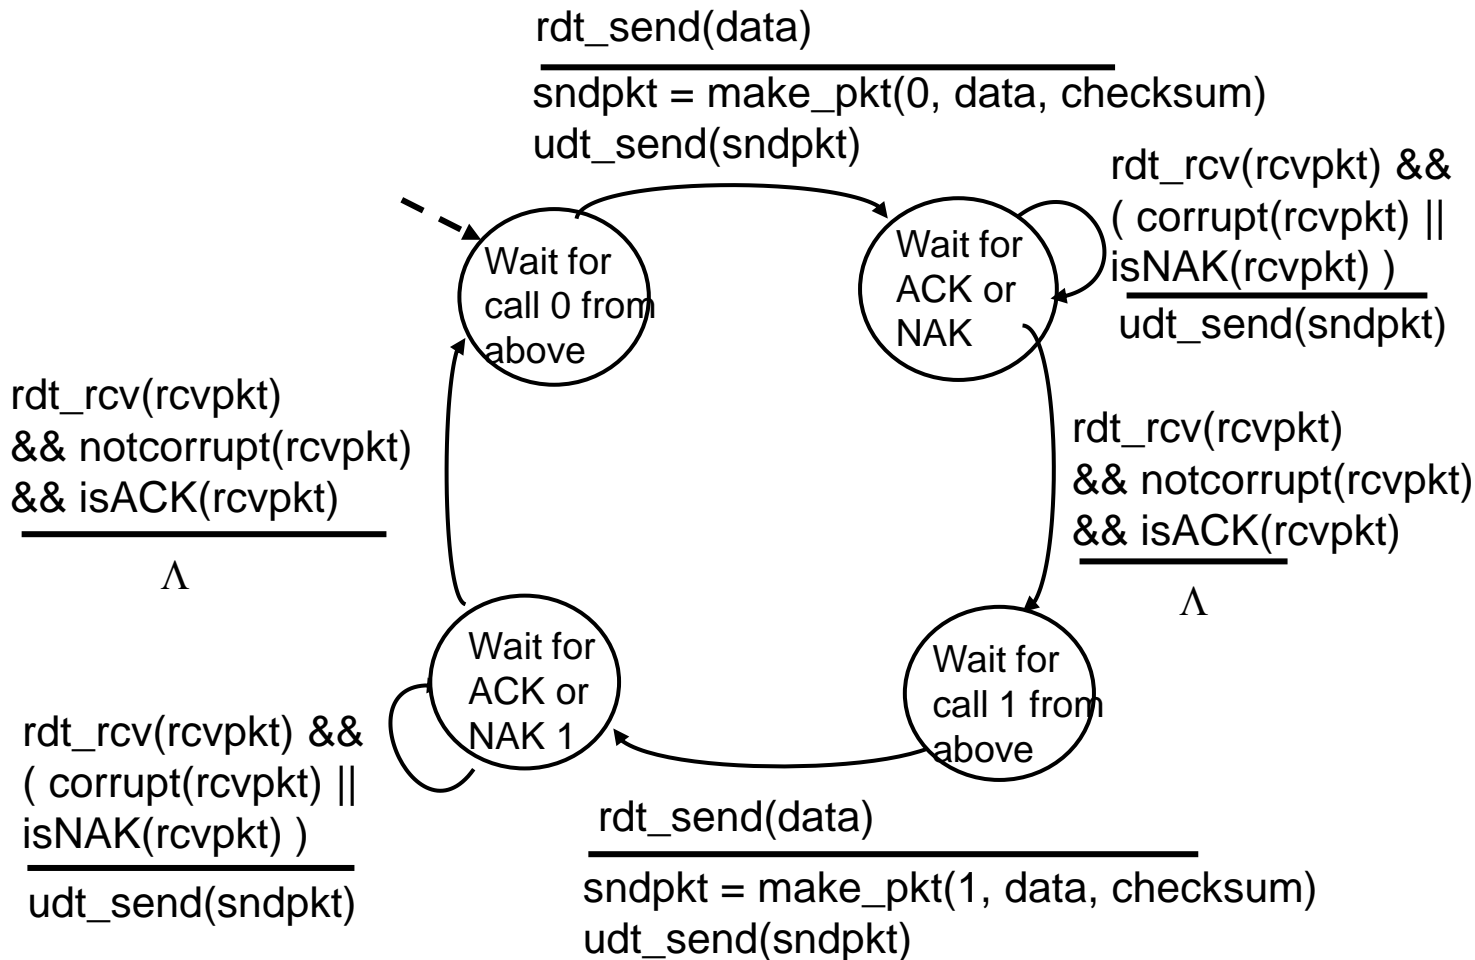


Execution traces of rdt2.1b are all that can be generated by the finite state machine above.

Recap: Protocol Analysis using State Invariants

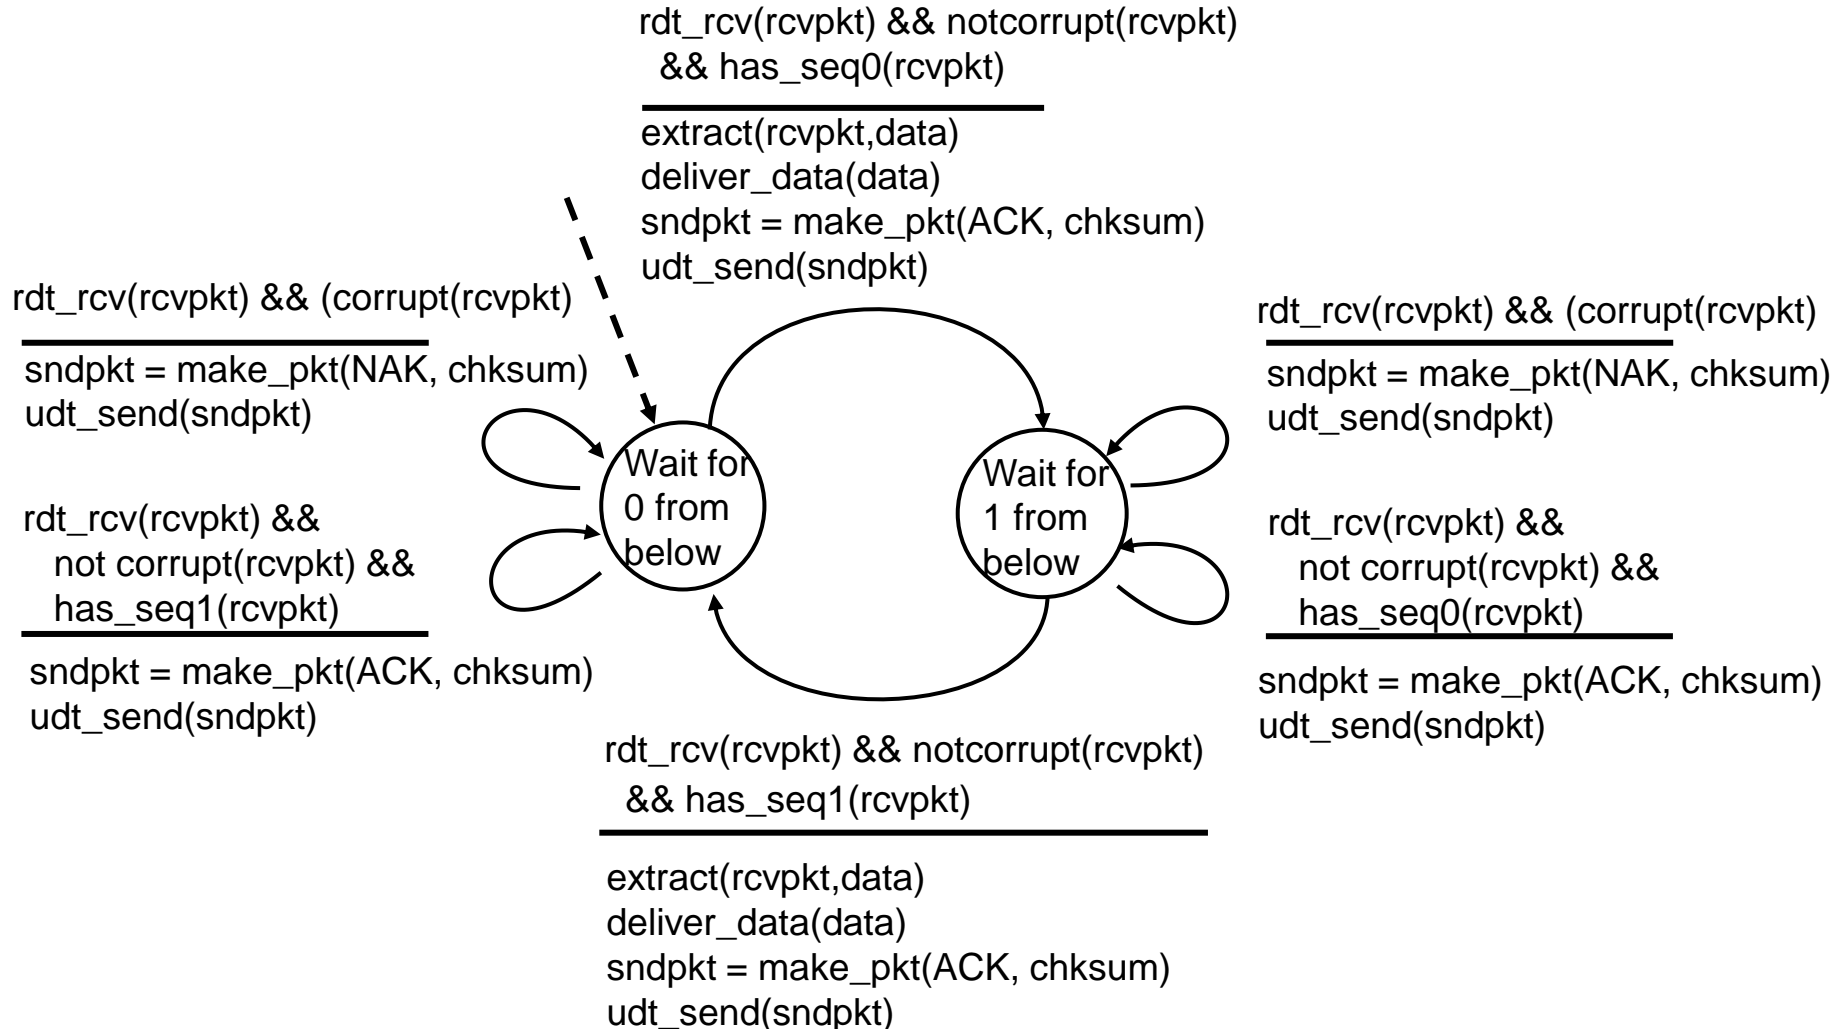


rdt2.1c: Sender, Handles Garbled ACK/NAKs: Using 1 bit (Alternating-Bit Protocol)



rdt2.1c: Receiver, Handles Garbled

ACK/NAKs: Using 1 bit



rdt2.1c: Summary

Sender:

- ❑ state must “remember” whether “current” pkt has 0 or 1 seq. #

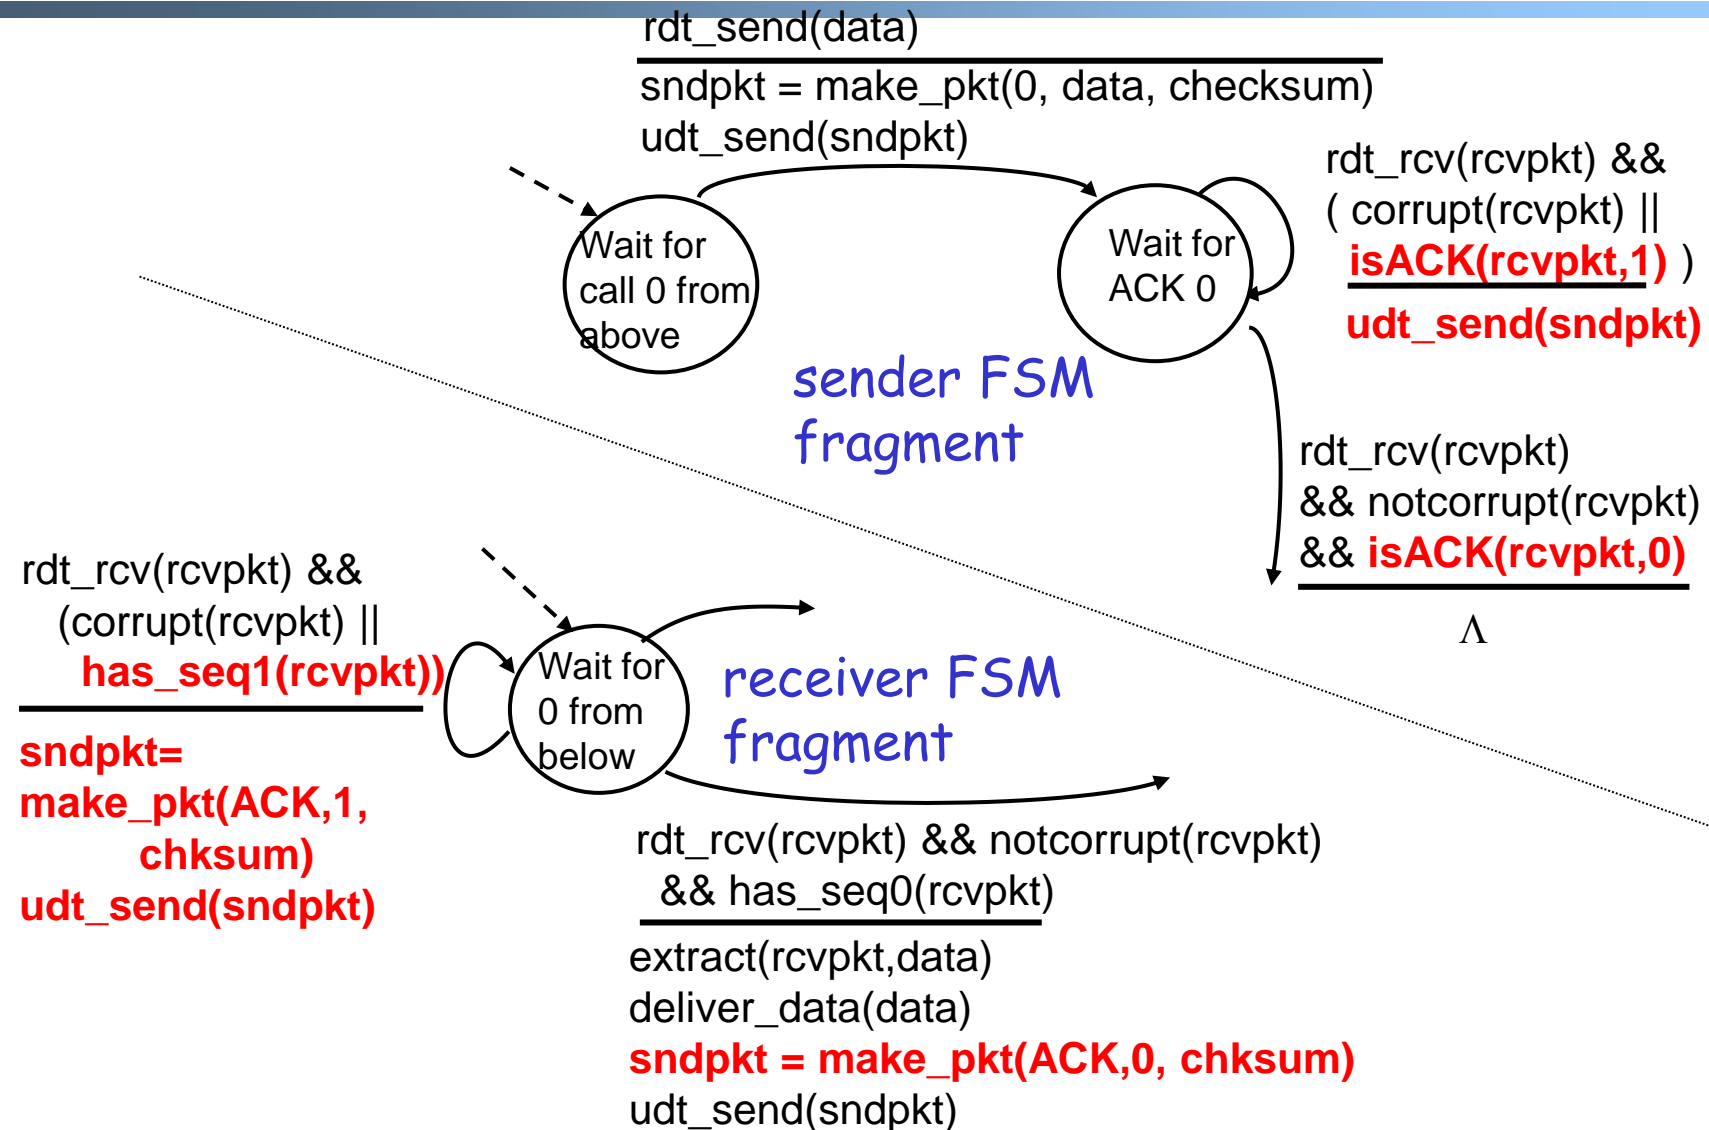
Receiver:

- ❑ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #

rdt2.2: a NAK-free protocol

- ❑ Same functionality as rdt2.1c, using ACKs only
- ❑ Instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- ❑ Duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: Sender, Receiver Fragments



Outline

- Admin and review
- Reliable data transfer
 - perfect channel
 - channel with bit errors
 - *channel with bit errors and losses*

rdt3.0: Channels with Errors and Loss

New assumption:

underlying channel can also lose packets (data or ACKs)

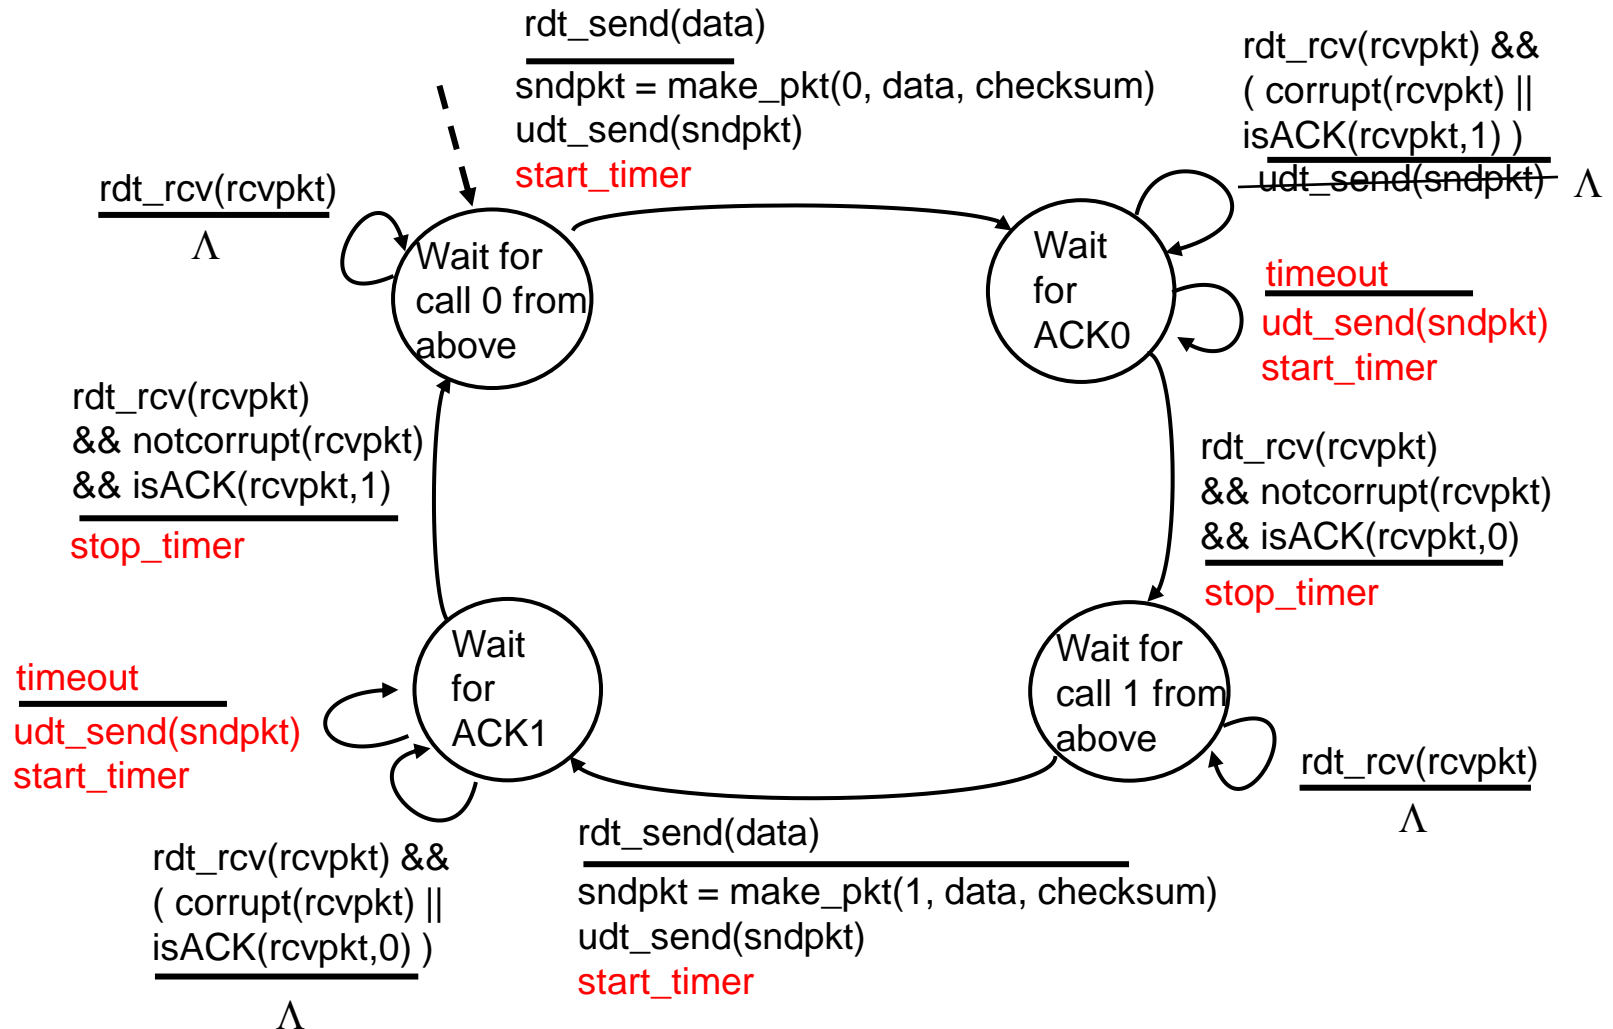
- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

Q: Does rdt2.2 work under losses?

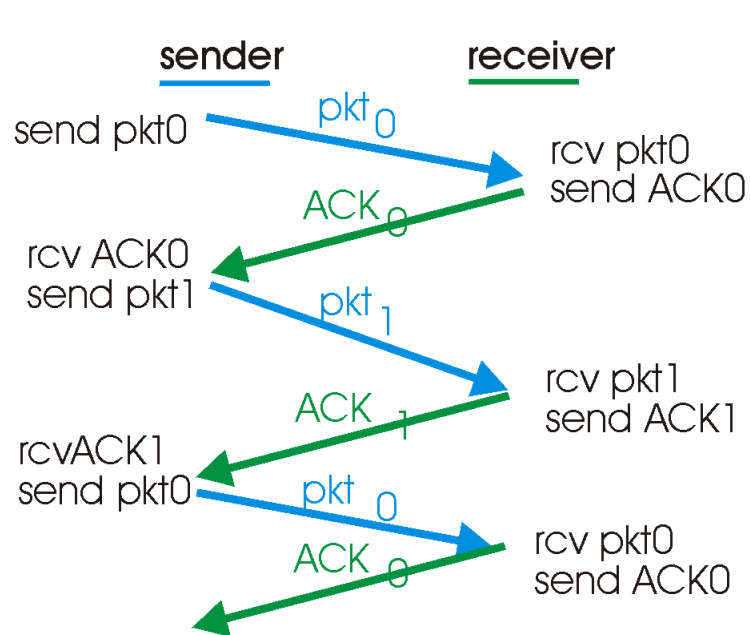
Approach: sender waits “reasonable” amount of time for ACK

- ❑ requires countdown timer
- ❑ retransmits if no ACK received in this time
- ❑ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed

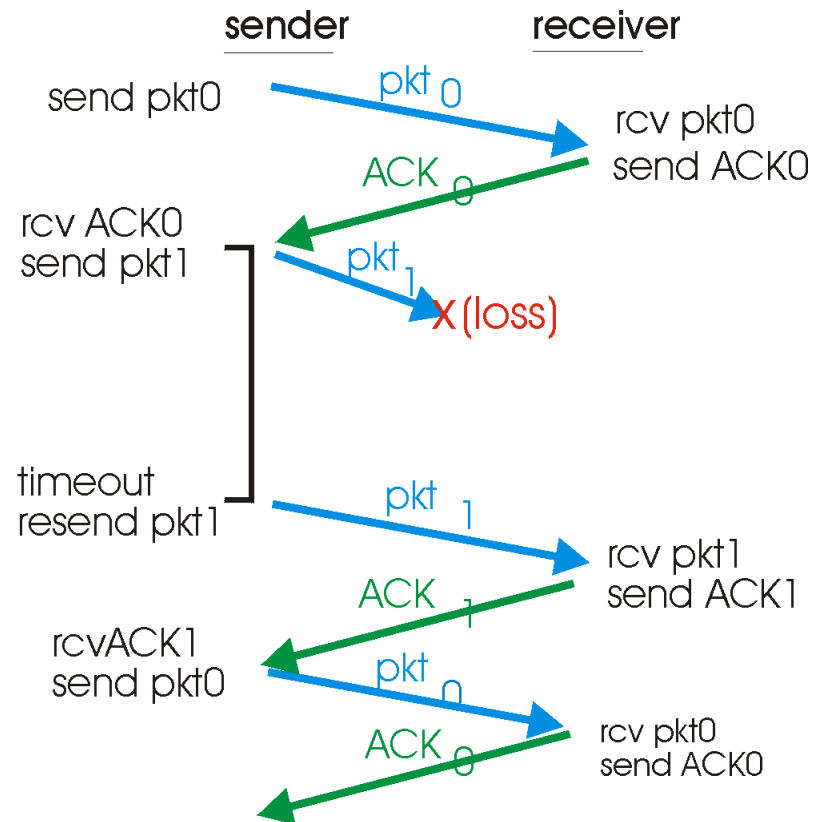
rdt3.0 Sender



rdt3.0 in Action




(a) operation with no loss

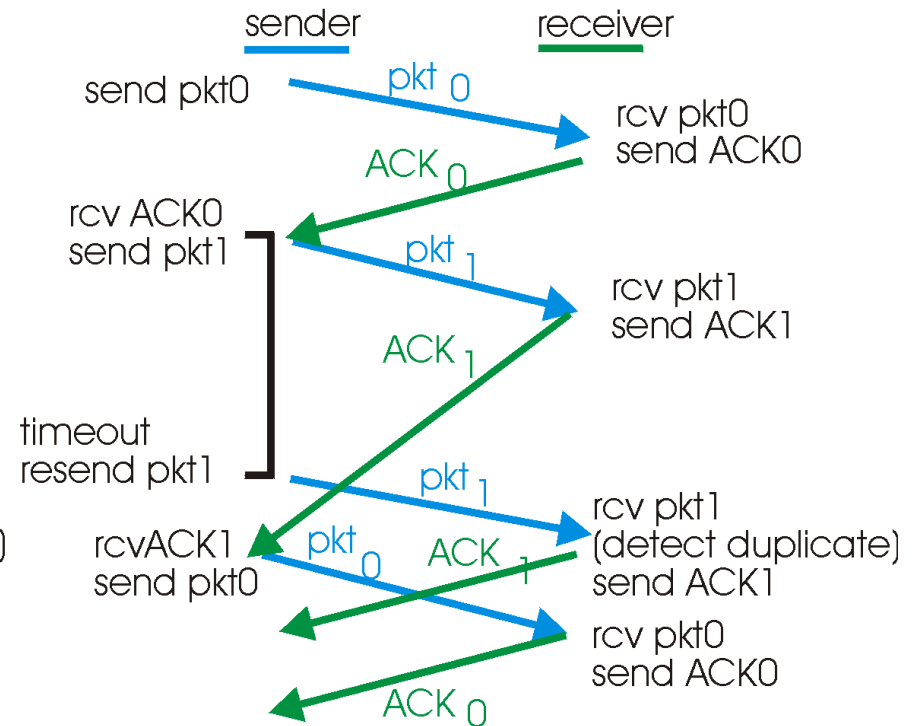


(b) lost packet

rdt3.0 in Action



(c) lost ACK




(d) premature timeout

Question to think about: How to determine a good timeout value?

Home exercise: What are execution traces of rdt3.0? What are some state invariants of rdt3.0?

rdt3.0: Stop-and-Wait Performance



What is U_{sender} : **utilization** – fraction of time link busy sending?

Assume: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet

Performance of rdt3.0

- ❑ rdt3.0 works, but performance stinks
- ❑ Example: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet:

$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^{**9} \text{ b/sec}} = 8 \text{ microsec}$$

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec -> 33kB/sec thruput over 1 Gbps link
- network protocol limits use of physical resources !


A Summary of Questions

- ❑ How to improve the performance of rdt3.0?
- ❑ What if there are reordering and duplication?
- ❑ How to determine the “right” timeout value?

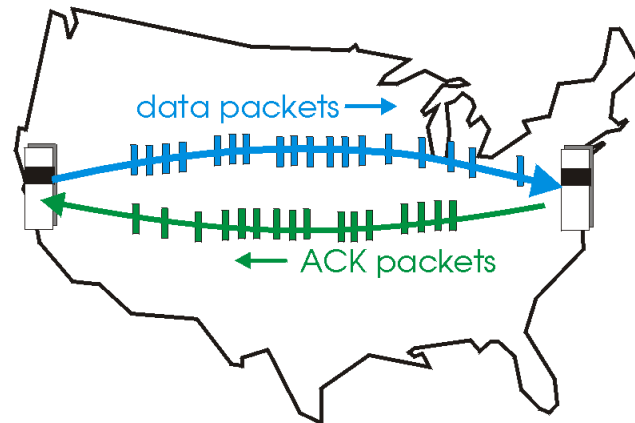
Sliding Window Protocols: Pipelining

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver

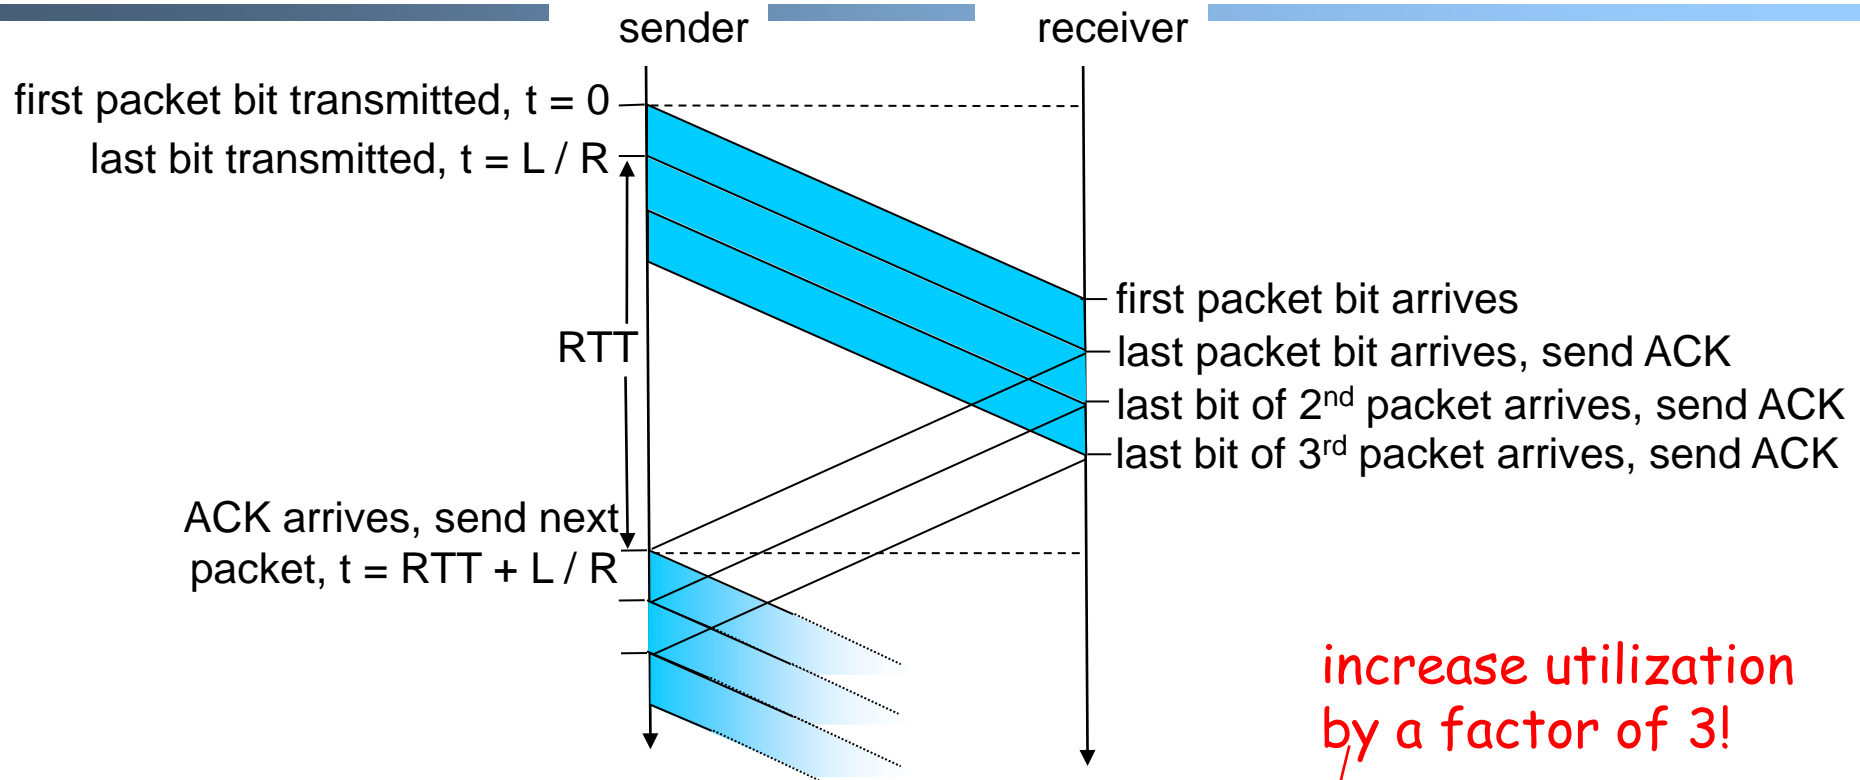


(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation

Pipelining: Increased Utilization



increase utilization
by a factor of 3!


$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Question: a rule-of-thumb window size?

Realizing Sliding Window: Go-Back-n


Sender:

- ❑ k-bit seq # in pkt header
- ❑ “window” of up to W , consecutive unack’ed pkts allowed

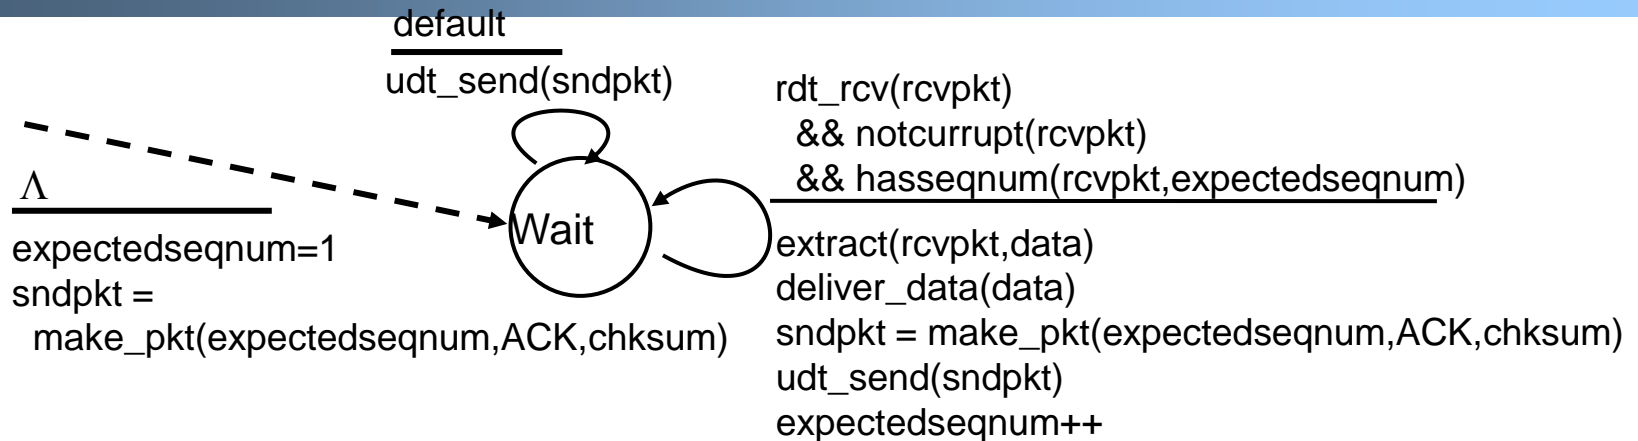


- ❑ **ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”**
 - note: ACK(n) could mean two things: I have received **upto and include** n , or I am waiting for n
- ❑ timer for the packet at base
- ❑ **timeout(n):** retransmit pkt n and all higher seq # pkts in window

GBN: Sender FSM



GBN: Receiver FSM



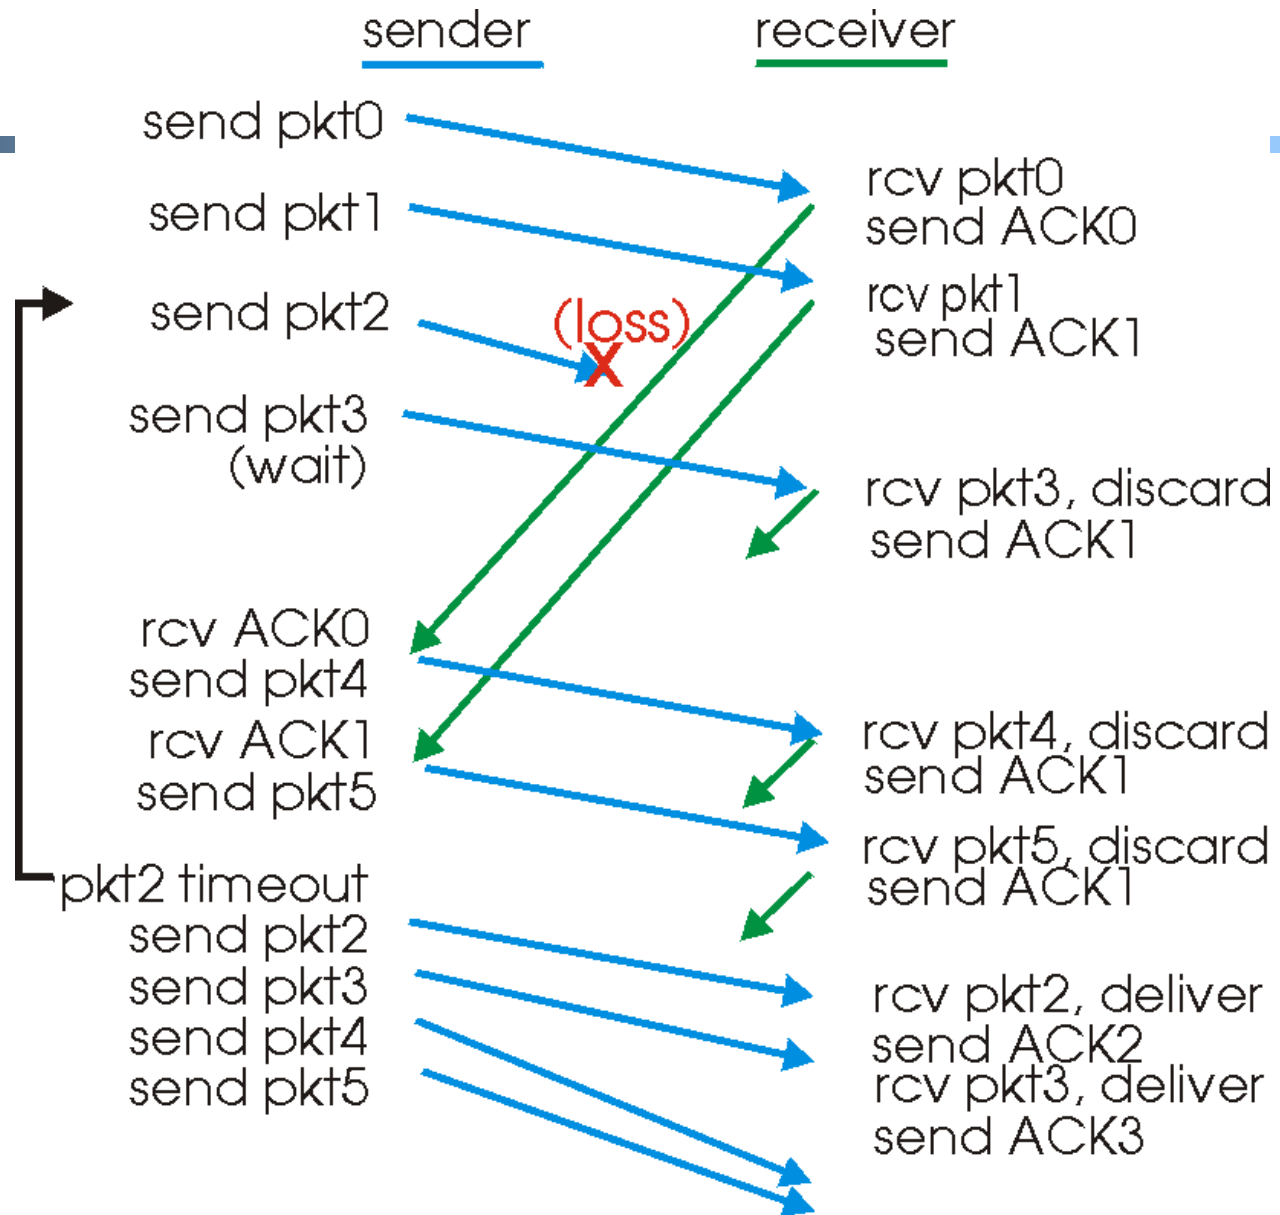
Only state: **expectedseqnum**

□ out-of-order pkt:

- discard (don't buffer) -> **no receiver buffering!**
- re-ACK pkt with highest in-order seq #
- may generate duplicate ACKs

GBN in Action

window
size = 4



Analysis: Efficiency of Go-Back-n

- ❑ Assume window size W
- ❑ Assume each packet is lost with probability p
- ❑ On average, how many packets do we send for each data packet received?

Selective Repeat

- ❑ Sender window
 - Window size W : W consecutive unACKed seq #'s
- ❑ Receiver *individually* acknowledges correctly received pkts
 - *buffers out-of-order* pkts, for eventual in-order delivery to upper layer
 - *ACK(n) means received packet with seq# n only*
 - buffer size at receiver: window size
- ❑ Sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt