

# Network Transport Layer: TCP

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

11/16/2021


# Admin

---

- ❑ Lab assignment 2 to be returned on Nov. 18
- ❑ Class Project
  - 15% of your final score
  - Please start ASAP
  - Talk to the instructor or TA to get feedback

# Recap: Reliable Transport

## □ Basic structure: sliding window protocols



(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation


General  
technique:  
pipelining.

## □ Realization: GBN or SR

# Recap: Go-Back-N

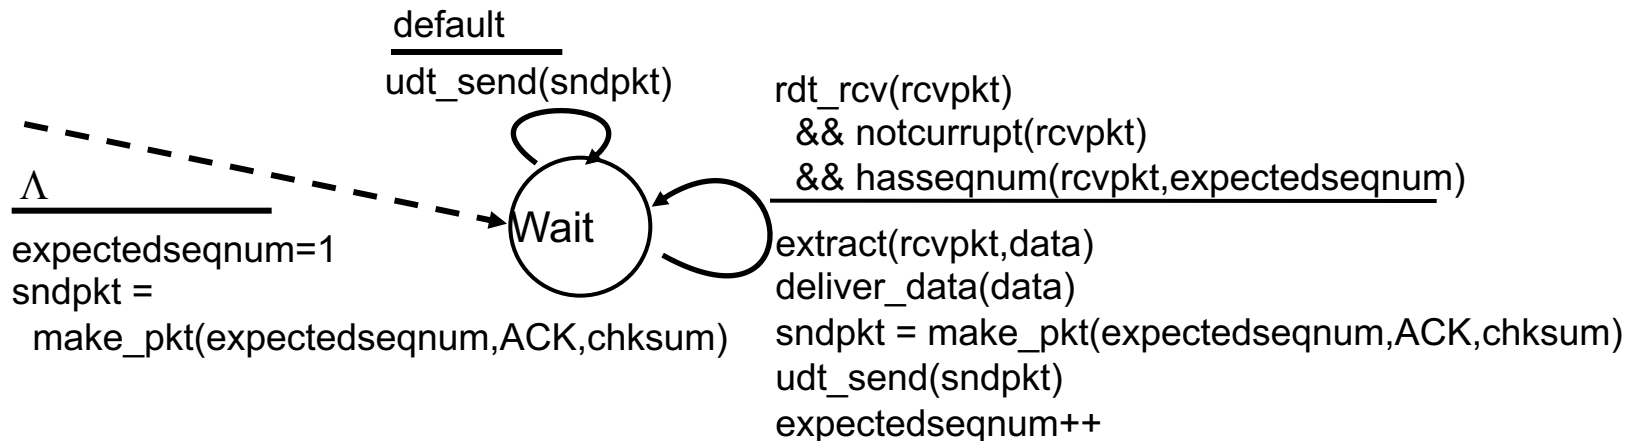
## Sender:

- ❑ k-bit seq # in pkt header
- ❑ “window” of up to  $W$ , consecutive unack’ed pkts allowed



- ❑ **ACK(n): ACKs all pkts up to, including seq #  $n$  - “cumulative ACK”**
  - note: ACK(n) could mean two things: I have received **upto and include**  $n$ , or I am waiting for  $n$
- ❑ timer for the packet at base
- ❑ *timeout(n)*: retransmit pkt  $n$  and all higher seq # pkts in window

# Recap: Go-Back-N



Only state: **expectedseqnum**

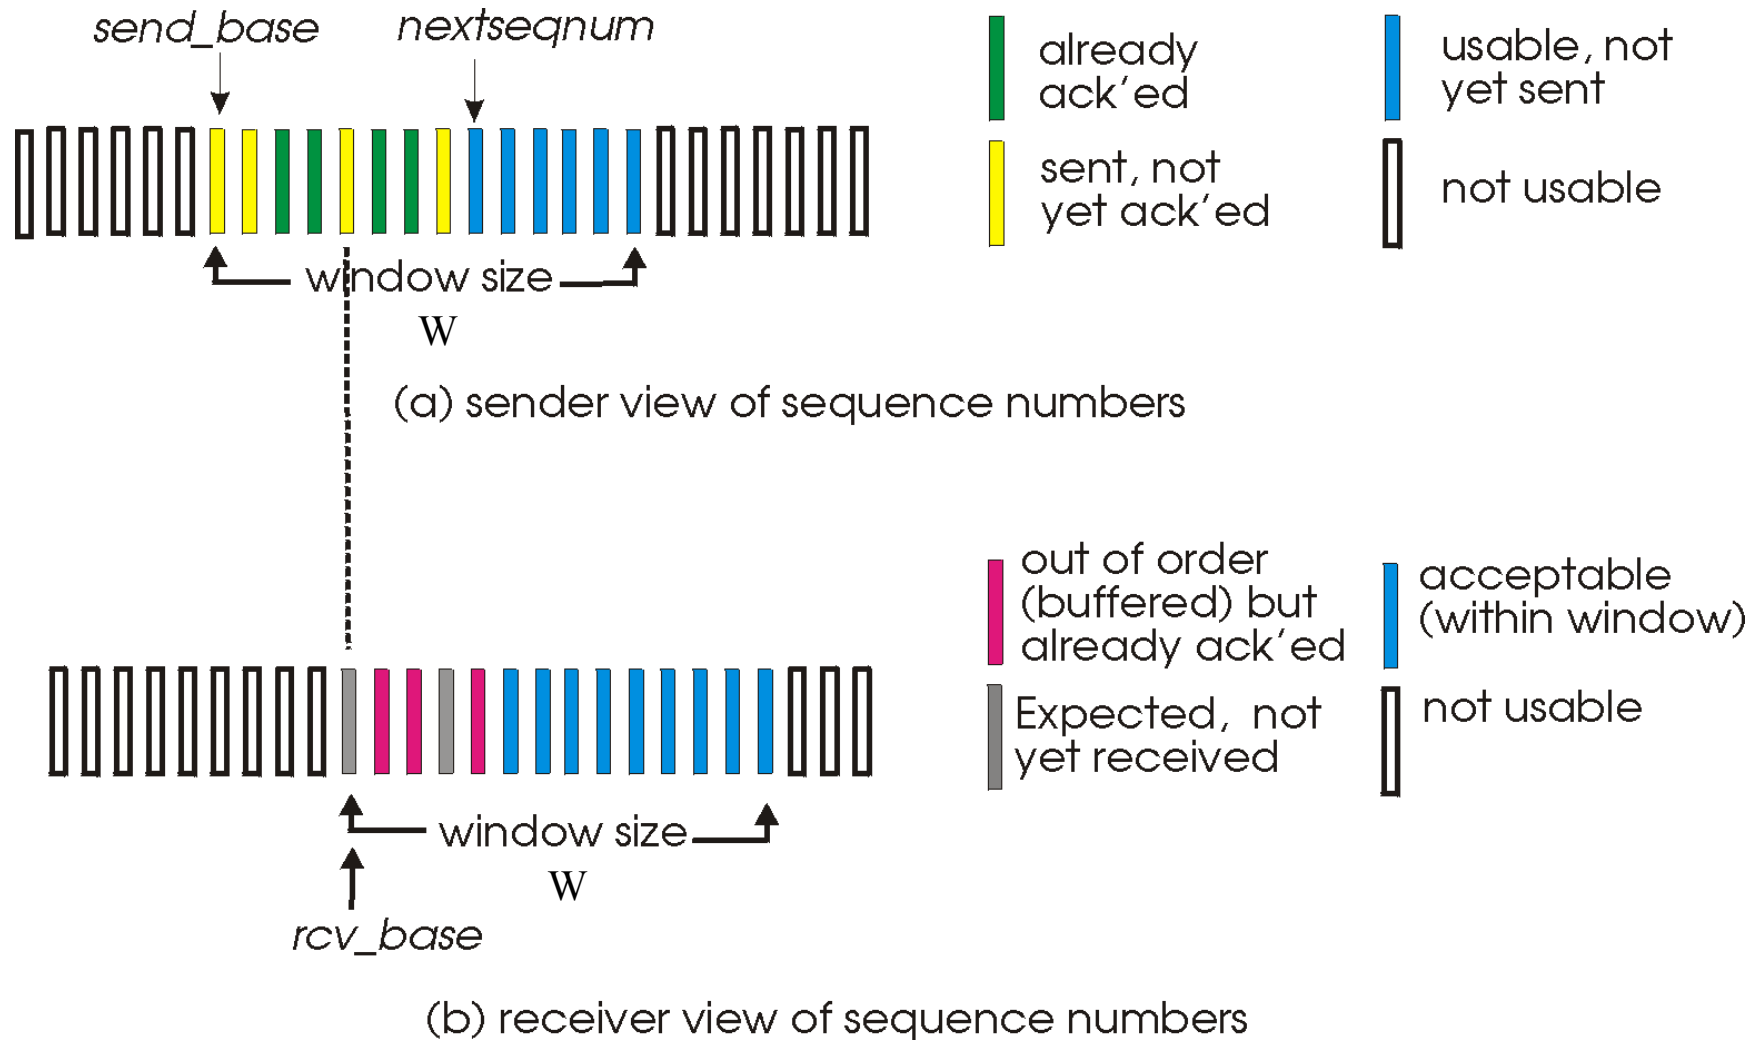
□ out-of-order pkt:

- discard (don't buffer) -> **no receiver buffering!**
- re-ACK pkt with highest in-order seq #
- may generate duplicate ACKs

# Selective Repeat

- ❑ Sender window
  - Window size  $W$ :  $W$  consecutive unACKed seq #'s
- ❑ Receiver *individually* acknowledges correctly received pkts
  - *buffers out-of-order* pkts, for eventual in-order delivery to upper layer
  - *ACK(n) means received packet with seq# n only*
  - buffer size at receiver: window size
- ❑ Sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt

# Selective Repeat: Sender, Receiver Windows



# Selective Repeat

## —sender—

data from above :

- unACKed packets is less than window size  $W$ , send; otherwise block app.

timeout(n):

- resend pkt  $n$ , restart timer

ACK(n) in [sendbase, sendbase+W-1]:

- mark pkt  $n$  as received
- update sendbase to the first packet unACKed

## —receiver—

pkt  $n$  in [rcvbase, rcvbase+W-1]

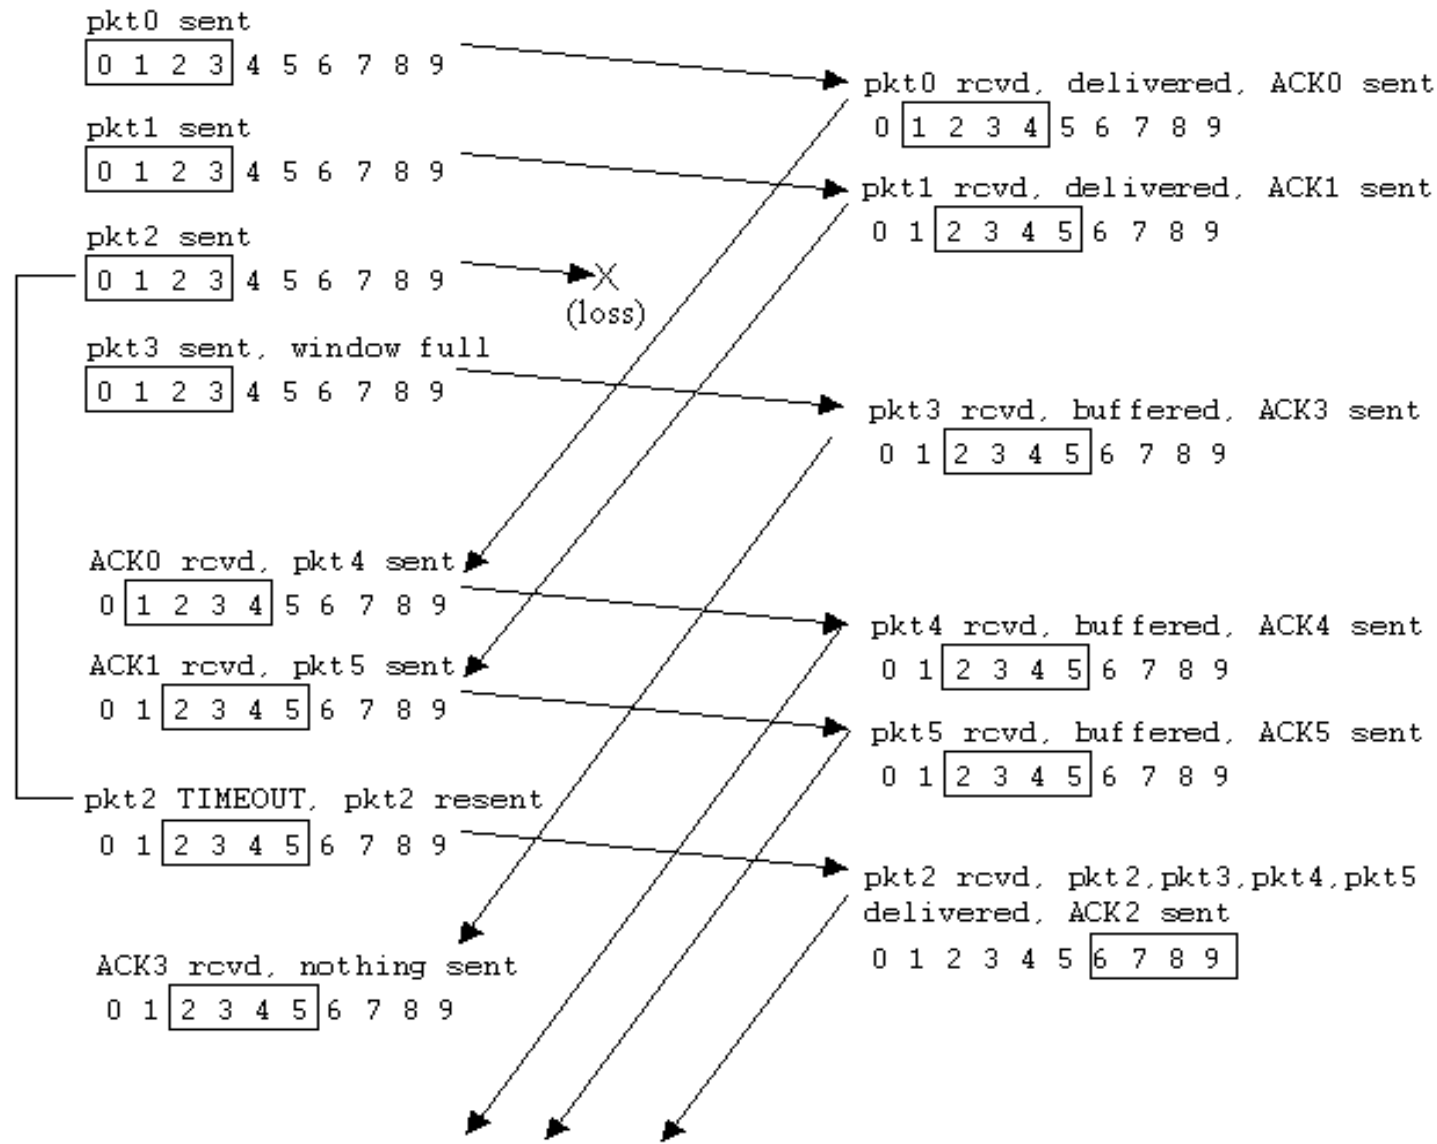
- send ACK(n)
- if (out-of-order)  
mark and buffer pkt  $n$   
else /\*in-order\*/  
deliver any in-order packets

otherwise:

- ignore



# Selective Repeat in Action




## Discussion: Efficiency of Selective Repeat

- ❑ Assume window size  $W$
- ❑ Assume each packet is lost with probability  $p$
- ❑ On average, how many packets do we send for each data packet received?

# Selective Repeat: Seq# Ambiguity


Example:

- ❑ seq #'s: 0, 1, 2, 3
- ❑ window size=3
- ❑ Error: incorrectly passes duplicate data as new.




# State Invariant: Window Location

## □ Go-back-n (GBN)




## □ Selective repeat (SR)




# Window Location

Q: what relationship between seq # size and window size?

## □ Go-back-n (GBN)



## □ Selective repeat (SR)



# Selective Repeat

## —sender—

data from above :

- unACKed packets is less than window size  $W$ , send; otherwise block app.

timeout(n):

- resend pkt  $n$ , restart timer

ACK(n) in [sendbase, sendbase+W-1]:

- mark pkt  $n$  as received
- update sendbase to the first packet unACKed

## —receiver—

pkt  $n$  in [rcvbase, rcvbase+W-1]

- send ACK(n)
- if (out-of-order)  
mark and buffer pkt  $n$   
else /\*in-order\*/  
deliver any in-order packets

pkt  $n$  in [rcvbase-W, rcvbase-1]

- send ACK(n)

otherwise:

- ignore

# Sliding Window Protocols: Go-back-n and Selective Repeat

	Go-back-n	Selective Repeat
data bandwidth: sender to receiver (avg. number of times a pkt is transmitted)	Less efficient $\frac{1-p+pw}{1-p}$	More efficient $\frac{1}{1-p}$
ACK bandwidth (receiver to sender)	More efficient	Less efficient
Relationship between M (the number of seq#) and W (window size)	$M > W$	$M \geq 2W$
Buffer size at receiver	1	W
Complexity	Simpler	More complex

p: the loss rate of a packet; M: number of seq# (e.g., 3 bit M = 8); W: window size

# Outline

---

- ❑ Admin and Recap
- ❑ Reliable data transfer
  - perfect channel
  - channel with bit errors
  - channel with bit errors and losses
  - sliding window: reliability with throughput
- *TCP reliability*



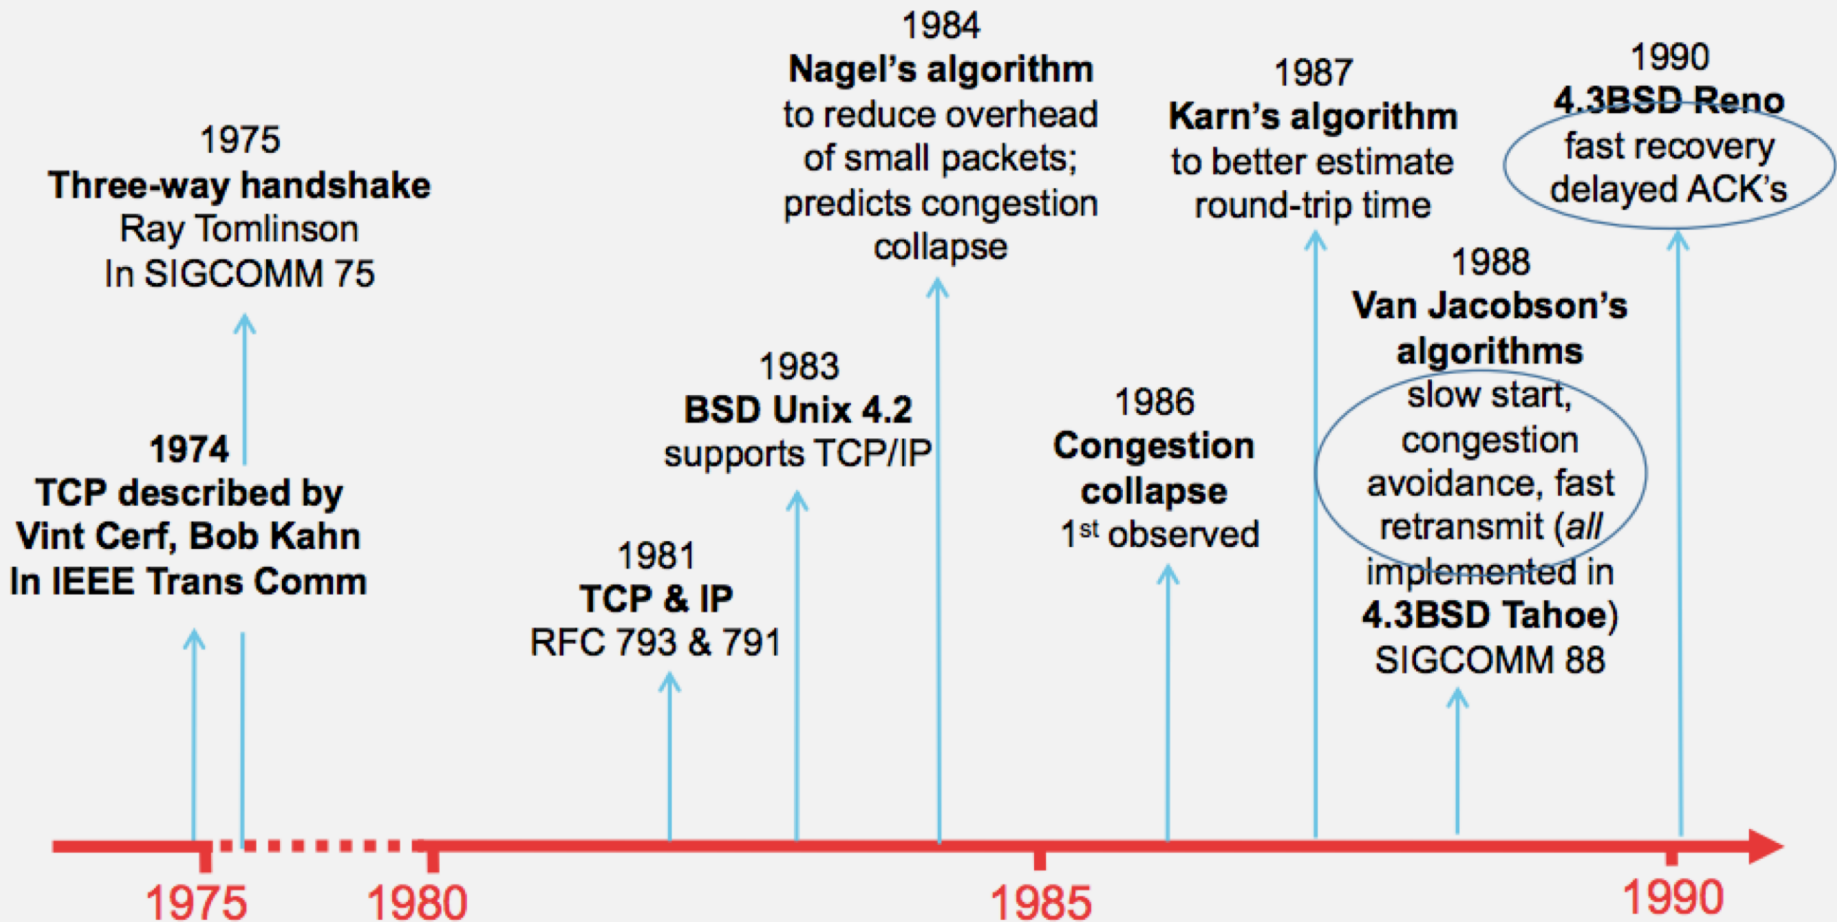
# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

---

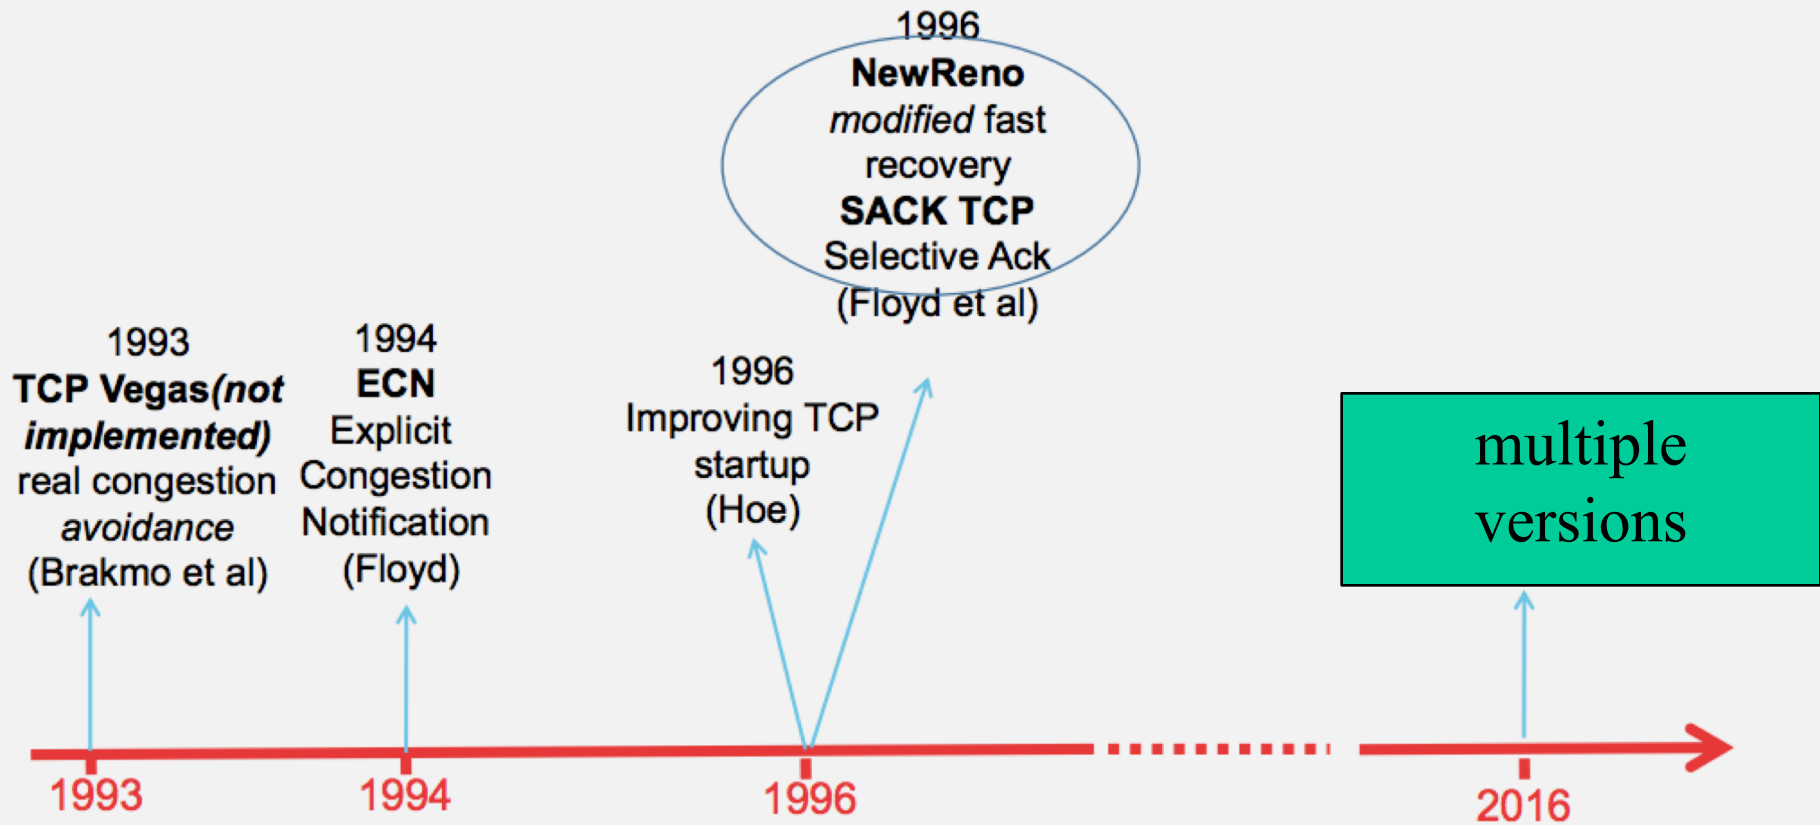
- Point-to-point reliability: one sender, one receiver
- Flow controlled and congestion controlled

# Evolution of TCP



Source: <http://webcourse.cs.technion.ac.il/236341/Winter2015-2016/ho/WCFiles/Tutorial10.pdf>

# Evolution of TCP



# TCP Reliable Data Transfer

## ❑ Connection-oriented:


- connection management
  - setup (exchange of control msgs) init's sender, receiver state before data exchange
  - close

## ❑ Full duplex data:

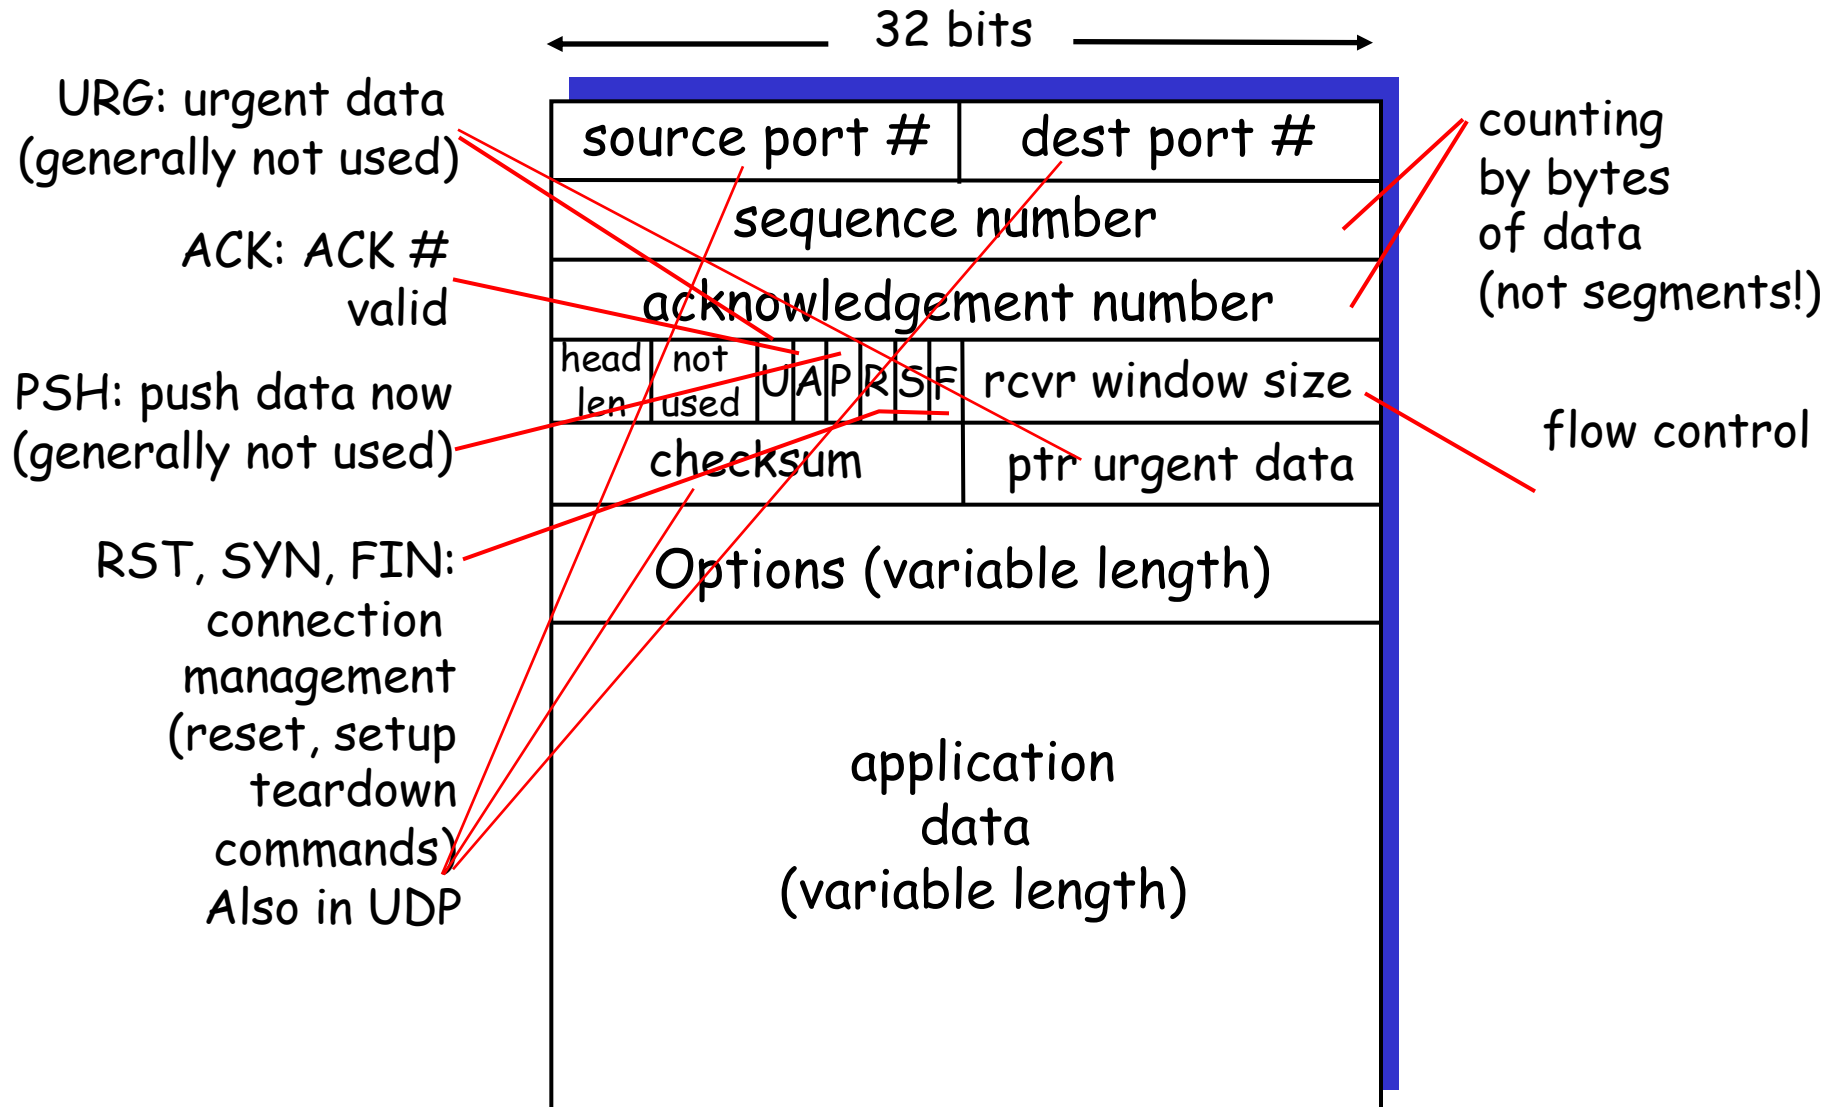
- bi-directional data flow in same connection

## ❑ A sliding window protocol

- a combination of go-back-n and selective repeat:
  - send & receive buffers
  - cumulative acks
  - TCP uses a single retransmission timer
  - do not retransmit all packets upon timeout



# TCP Segment Structure

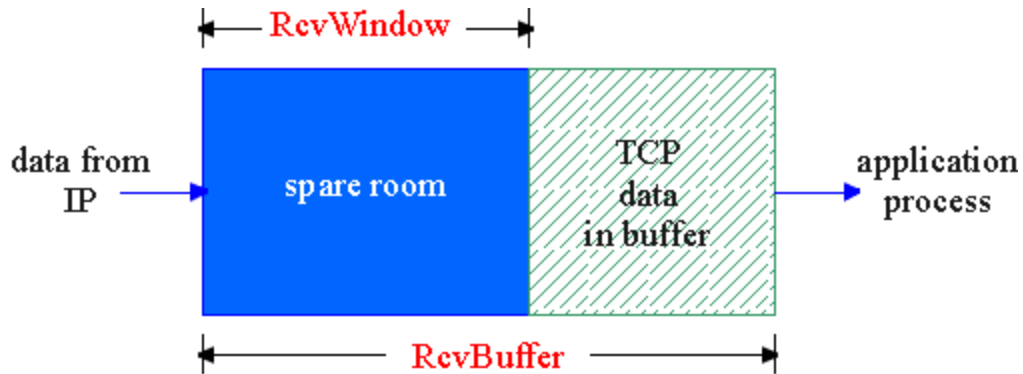


# Outline

- ❑ Admin and Recap
- ❑ Reliable data transfer
  - perfect channel
  - channel with bit errors
  - channel with bit errors and losses
  - sliding window: reliability with throughput
- ❑ TCP reliability
  - *data seq#, ack, buffering*

# Flow Control

- ❑ receive side of a connection has a receive buffer:




## flow control

sender won't overflow receiver's buffer by transmitting too much, too fast

- ❑ speed-matching service: matching the send rate to the receiving app's drain rate

- ❑ app process may be slow at reading from buffer

# TCP Flow Control: How it Works



□ spare room in buffer  
= **RcvWindow**

source port #					dest port #				
sequence number									
acknowledgement number									
head len	not used	U	A	P	R	S	F	rcvr window size	
checksum					ptr urgent data				
Options (variable length)									
application data (variable length)									



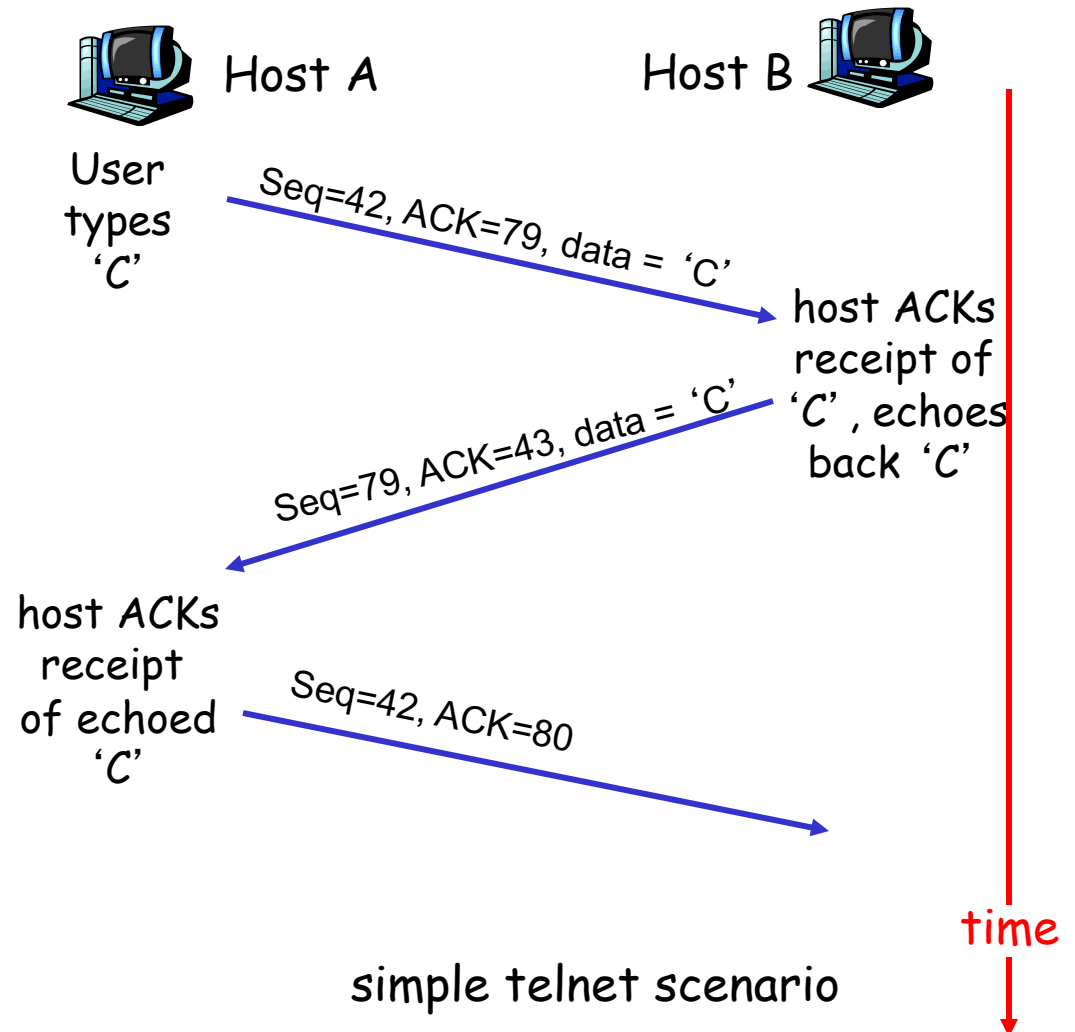
# TCP Seq. #'s and ACKs

## Seq. #'s:

- byte stream  
“number” of first  
byte in segment's  
data

## ACKs:

- seq # of next byte  
**expected** from  
other side
- **cumulative** ACK in  
standard header
- selective ACK in  
options



# TCP Send/Ack Optimizations

- ❑ TCP includes many tune/optimizations, e.g.,
  - the "small-packet problem": sender sends a lot of small packets (e.g., telnet one char at a time)
    - Nagle's algorithm: do not send data if there is small amount of data in send buffer and there is an unack'd segment
  - the "ack inefficiency" problem: receiver sends too many ACKs, no chance of combining ACK with data
    - Delayed ack to reduce # of ACKs/combine ACK with reply

# TCP Receiver ACK Generation [RFC 1122, RFC 2581]

Event at Receiver	TCP Receiver Action
Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
Arrival of in-order segment with expected seq #. One other segment has ACK pending	Immediately send single cumulative ACK, ACKing both in-order segments
Arrival of out-of-order segment higher-than-expect seq. # . Gap detected	Immediately send duplicate ACK, indicating seq. # of next expected byte
Arrival of segment that partially or completely fills gap	Immediate send ACK, provided that segment starts at lower end of gap

# Outline

- ❑ Admin and Recap
- ❑ Reliable data transfer
  - perfect channel
  - channel with bit errors
  - channel with bit errors and losses
  - sliding window: reliability with throughput
- ❑ TCP reliability
  - data seq#, ack, buffering
  - *timeout realization*

# TCP Reliable Data Transfer

- ❑ Basic structure: sliding window protocol
- ❑ Remaining issue: How to determine the “right” parameters?
  - timeout value?
  - sliding window size?

# History

- ❑ Key parameters for TCP in mid-1980s
  - fixed window size  $W$
  - timeout value = 2 RTT
- ❑ Network collapse in the mid-1980s
  - UCB  $\leftrightarrow$  LBL throughput dropped by 1000X !
- ❑ The intuition was that the collapse was caused by wrong parameters...

# Timeout: Cost of Timeout Param

Why is good timeout value important?

❑ too short

- premature timeout
- unnecessary retransmissions; many duplicates

❑ too long

- slow reaction to segment loss


Q: Is it possible to set Timeout as a constant?

Q: Any problem w/ the early approach:  $\text{Timeout} = 2 \text{ RTT}$

# Setting Timeout

## Problem:

- ❑ Ideally, we set timeout = RTT, but RTT is not a fixed value  
=> using the average of RTT will generate many timeouts due to network variations
- ❑ Possibility: using the average/median of RTT
- ❑ Issue: this will generate many timeouts due to network variations



## Solution:

- ❑ Set Timeout RTO = avg + “safety margin” based on variation

TCP approach:

$$\text{Timeout} = \text{EstRTT} + 4 * \text{DevRTT}$$




# Compute EstRTT and DevRTT

- ❑ Exponential weighted moving average (EWMA)
  - influence of past sample decreases exponentially fast

$$\text{EstRTT} = (1-\alpha) * \text{EstRTT} + \alpha * \text{SampleRTT}$$


- **SampleRTT**: measured time from segment transmission until ACK receipt
- typical value:  $\alpha = 0.125$



$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstRTT}|$$

(typically,  $\beta = 0.25$ )

# An Example TCP Session



# Fast Retransmit

- ❑ Issue: Timeout period often relatively long:
  - long delay before resending lost packet
- ❑ Question: Can we detect loss faster than RTT?
  
- ❑ Detect lost segments via duplicate ACKs
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs
- ❑ If sender receives 3 ACKs for the same data, it supposes that segment after ACKed data was lost:
  - resend segment before timer expires

# Triple Duplicate Ack

Packets



Acknowledgements (waiting seq#)



# Fast Retransmit:

```
event: ACK received, with ACK field value of y
    if (y > SendBase) {
        ...
        SendBase = y
        if (there are currently not-yet-acknowledged segments)
            start timer
        ...
    }
    else {
        increment count of dup ACKs received for y
        if (count of dup ACKs received for y = 3) {
            resend segment with sequence number y
        }
        ...
    }
```

a duplicate ACK for  
already ACKed segment

fast retransmit

# TCP: reliable data transfer


Simplified  
TCP  
sender

```
00 sendbase = initial_sequence number agreed by TWH
01 nextseqnum = initial_sequence number by TWH
02 loop (forever) {
03     switch(event)
04     event: data received from application above
05         if (window allows send)
06             create TCP segment with sequence number nextseqnum
06             if (no timer) start timer
07             pass segment to IP
08             nextseqnum = nextseqnum + length(data)
           else put packet in buffer
09     event: timer timeout for sendbase
10         retransmit segment
11         compute new timeout interval
12         restart timer
13     event: ACK received, with ACK field value of y
14         if (y > sendbase) { /* cumulative ACK of all data up to y */
15             cancel the timer for sendbase
16             sendbase = y
17             if (no timer and packet pending) start timer for new sendbase
17             while (there are segments and window allow)
18                 sent a segment;
18         }
19         else { /* y==sendbase, duplicate ACK for already ACKed segment */
20             increment number of duplicate ACKs received for y
21             if (number of duplicate ACKS received for y == 3) {
22                 /* TCP fast retransmit */
23                 resend segment with sequence number y
24                 restart timer for segment y
25             }
26     } /* end of loop forever */
```

# Outline


- ❑ Admin and Recap
- ❑ Reliable data transfer
  - perfect channel
  - channel with bit errors
  - channel with bit errors and losses
  - sliding window: reliability with throughput
- ❑ TCP reliability
  - data seq#, ack, buffering
  - timeout realization
  - *connection management*

# Why Connection Setup/When to Accept (Safely Deliver) First Packet?





# Why Connection Setup/When to Accept (Safely Deliver) First Packet?



# Transport "Safe-Setup" Principle

- A general safety principle for a receiver R to accept a message from a sender S is the general "**authentication**" principle, which consists of two conditions:

Transport authentication principle:

- [p1] Receiver can be sure that what Sender says is **fresh**
- [p2] Receiver receives something that **only Sender can say**

We first assume a secure setting: no malicious attacks.

Exercise: Techniques to allow a receiver to check for freshness (e.g., add a time stamp)?